

PONDICHERRY UNIVERSITY
RAMANUJAN SCHOOL OF MATHEMATICAL SCIENCES
DEPARTMENT OF STATISTICS



M.Sc. QUANTITATIVE FINANCE
SEMESTER III
LECTURE NOTES ON
Machine Learning in Quantitative Finance

PREPARED BY
NAME: N Rohit Vedhanandh
REG. NO.: 25MSQUFPY0002

OCTOBER 2026

Contents

1	Machine Learning Foundations & Data Preprocessing	1
1.1	Challenges in Machine Learning	1
1.2	Types of Variables	1
1.3	Data Preprocessing Pipeline	1
1.3.1	Scaling & Normalization	2
1.4	Train-Test-Validation Splitting	3
1.4.1	Common Partitioning Ratios	3
1.4.2	K-Fold Cross-Validation	3
1.4.3	Data Leakage Prevention	3
1.4.4	Splitting Financial Time-Series Data	3
1.5	Missing Data Mechanisms & Imputation	4
1.5.1	Missingness Mechanisms	4
1.5.2	Imputation Methods	4
1.5.3	Decision Framework for Selecting Imputation Methods	4
1.6	Imbalanced Categorical Data	4
1.6.1	Consequences of Class Imbalance	5
1.6.2	Remedial Solutions	5
1.7	Model Evaluation Metrics	5
1.8	Mistakes to be Avoided in ML Pipelines	6
1.9	Key Takeaways	6
2	K-Nearest Neighbors (KNN)	8
2.1	Introduction to KNN	8
2.1.1	Core Algorithmic Characteristics	8
2.1.2	Algorithmic Workflow	8
2.2	Distance Metrics	8
2.2.1	Euclidean Distance	9
2.2.2	Manhattan Distance	9
2.2.3	Minkowski Distance	9
2.3	Worked Numerical Examples	9
2.3.1	Fruit Classification via Euclidean Distance	9
2.3.2	Fintech Stock Recommendation with Standardized Indicators	11
2.3.3	Bank Loan Approval via Manhattan Distance	12
2.3.4	Real Estate House Price Prediction (KNN Regression)	13
2.4	Parameter Selection: Small k vs. Large k	14
3	Naive Bayes Classification	15
3.1	Probability Foundations & Notations	15
3.1.1	Key Probability Identities	15
3.1.2	Conditional Probability & Multiplication Rule	15
3.1.3	Statistical Independence	16

3.2	Bayes' Theorem	16
3.3	Worked Numerical Example: Credit Risk Classification	16
3.4	The Bayes Classifier Framework	17
3.4.1	Marginal Evidence Across K Classes	17
3.4.2	Optimal Bayes Decision Rule	17
3.5	The Naive Bayes Classifier	17
3.5.1	The Challenge of High-Dimensional Likelihoods	17
3.5.2	The Naive Conditional Independence Assumption	18
3.5.3	Naive Bayes Classification Rule	18
3.5.4	Algorithmic Step-by-Step Sequence	18
4	Logistic Regression	20
4.1	Introduction and Core Motivation	20
4.1.1	Why Not Standard Linear Regression?	20
4.2	The Sigmoid (Logistic) Function and Probability Modeling	20
4.2.1	Algebraic Equivalence of the Sigmoid Function	21
4.2.2	Properties of the Sigmoid Curve and Decision Threshold	21
4.3	Key Terms and Matrix Notation	22
4.3.1	Matrix Representation	22
4.4	Odds, Log-Odds (Logit), and Parameter Interpretation	23
4.4.1	Definition of Odds	23
4.4.2	Log-Odds (Logit Function)	23
4.4.3	Interpreting Coefficients and Odds Ratios	24
4.5	Model Training via Maximum Likelihood Estimation (MLE)	24
4.6	Assumptions of Logistic Regression	24
4.7	Variants, Metrics, and Trade-Offs	25
4.7.1	Types of Logistic Regression	25
4.7.2	Model Evaluation Metrics	25
4.7.3	Advantages and Disadvantages	25
4.8	End-to-End Workflow and Generalized Linear Model (GLM) Classification	26
4.8.1	Five-Step Implementation Pipeline	26
4.8.2	Membership in Generalized Linear Models (GLM)	26
5	Decision Trees	27
5.1	Overview and Core Concepts	27
5.1.1	Structural Components of a Decision Tree	27
5.1.2	General Recursive Splitting Algorithm	27
5.2	Impurity Measures and Splitting Criteria	28
5.2.1	Standard Tree Induction Algorithms	28
5.2.2	Entropy and Information Gain	28
5.2.3	Gain Ratio (C4.5)	29
5.2.4	Gini Impurity (CART)	29
5.2.5	Comparison: Gini Impurity vs. Entropy	29
5.2.6	Advantages and Disadvantages of Decision Trees	29
5.3	Solved Problem 5.1: Play Tennis Decision Tree	30
5.4	Solved Problem 5.2: Credit Loan Approval Decision Tree	34
5.5	Solved Problem 5.3: Candidate Job Offer Decision Tree	38

6 Support Vector Machines	42
6.1 Core Principles and Mathematical Foundations	42
6.1.1 Fundamental Definitions and Geometry	42
6.1.2 Functional Margin and Hard-Margin Condition	43
6.2 Boundary Identification and Analytical Solution Techniques	43
6.2.1 Horizontal Separation ($x_2 = c$)	43
6.2.2 Vertical Separation ($x_1 = c$)	43
6.2.3 Diagonal Separation ($x_2 - x_1 = c$ or $x_2 + x_1 = c$)	44
6.3 Solved Problem 6.1: Horizontal Decision Boundary	44
6.4 Solved Problem 6.2: Vertical Decision Boundary	46
6.5 Solved Problem 6.3: Diagonal Decision Boundary	48
Comprehensive Formula Reference Sheet	51
Chapter 1: Machine Learning Foundations and Data Preprocessing	51
Chapter 2: K-Nearest Neighbors (KNN)	52
Chapter 3: Naive Bayes Classification	53
Chapter 4: Logistic Regression	54
Chapter 5: Decision Trees	56
Chapter 6: Support Vector Machines (SVM)	57

Chapter 1

Machine Learning Foundations & Data Preprocessing


1.1 Challenges in Machine Learning

Machine learning applications in financial engineering face several structural and regulatory constraints:

- **Requires Large, Clean Data:** Models demand substantial volumes of high-quality, cleansed historical records to generalize effectively.
- **Risk of Overfitting:** Complex models risk fitting noise and idiosyncratic market movements rather than genuine underlying signals.
- **Interpretability (“Black Box”):** Advanced algorithms often lack transparent mechanistic explanations, creating auditing hurdles.
- **Regulatory & Ethical Concerns in Finance:** Strict regulatory compliance mandates model explainability, fairness, and accountability in financial decision-making.

1.2 Types of Variables

Variables encountered in machine learning datasets are broadly categorized into numerical and categorical types:




1.3 Data Preprocessing Pipeline

The standard end-to-end data preparation workflow transforms raw financial inputs into model-ready datasets:

- **Cleaning:** Removes duplicate records, corrects erroneous entries, and handles missing observations.

Table 1.1: Taxonomy of Data Types and Financial Examples

Type	Subtype	Meaning	Finance Example
Numerical	Continuous	Any value in a continuous range	Stock price, portfolio return, volatility
Numerical	Discrete	Countable whole numbers	Number of trades per day, number of defaults
Categorical	Nominal	No natural ordering among categories	Sector: Banking, IT, Pharma
Categorical	Ordinal	Natural hierarchical order exists	Credit Rating: AAA > AA > A > BBB
Categorical	Binary	Exactly two distinct categories	Loan Default: Yes / No

**Figure 1.1:** End-to-End Data Preprocessing Pipeline from Raw Inputs to Model-Ready Data.

- **Transformation:** Converts raw figures into mathematically suitable modeling formats (e.g., calculating log returns from sequential stock prices).
- **Normalization / Scaling:** Brings numerical features across heterogeneous magnitudes into a common comparable scale.
- **Encoding:** Converts non-numeric categorical attributes into numerical indicator representations.

1.3.1 Scaling & Normalization

Features measured in disparate units must be rescaled to ensure distance-based and gradient-based algorithms treat features proportionately:

Min-Max Scaling

Rescales feature values into the bounded closed interval $[0, 1]$:

$$x' = \frac{x - x_{\min}}{x_{\max} - x_{\min}} \quad (1.1)$$

Z-Score Standardization

Standardizes observations using empirical sample mean and sample standard deviation:

$$Z = \frac{x - \mu}{\sigma} \quad (1.2)$$

1.4 Train-Test-Validation Splitting

Datasets are partitioned into distinct subsets to evaluate model performance on unseen data, verifying whether the algorithm generalizes rather than memorizing training samples:

- **Training Set:** Used to fit model parameters and learn patterns.
- **Validation Set:** Used during model development to tune hyperparameters and guide model selection.
- **Testing Set:** Used strictly for final unbiased model performance evaluation.

1.4.1 Common Partitioning Ratios

- **Train-Test Splits:** 80% – 20% or 70% – 30%.
- **Train-Validation-Test Splits:** 70% – 15% – 15%.

1.4.2 K-Fold Cross-Validation

In k -fold cross-validation:

1. The training dataset is partitioned into k equal parts (“folds”).
2. The model is trained on $k - 1$ folds and validated on the remaining fold.
3. This procedure is repeated k times, rotating the held-out validation fold sequentially across all splits.

Benefits & Common Choices

- Provides a significantly more robust estimate of true out-of-sample model performance.
- Reduces vulnerability to an anomalous single lucky or unlucky split.
- Common configurations: $k = 5$ or $k = 10$.

1.4.3 Data Leakage Prevention

Data Leakage Warning

Never scale or normalize the entire dataset prior to splitting.

If sample mean μ and standard deviation σ are calculated over the aggregate dataset, information from the test set leaks directly into the training stage.

Mandatory Pipeline Sequence:

1. Split the raw dataset first.
2. Fit the scaler/imputer exclusively on training data.
3. Transform validation and test data using the saved training parameters.

1.4.4 Splitting Financial Time-Series Data

For time-series data like stock prices, a purely random split is fundamentally problematic because it mixes past and future observations. Sequential financial records must strictly use **chronological splits**.

1.5 Missing Data Mechanisms & Imputation

1.5.1 Missingness Mechanisms

Table 1.2: Mechanisms of Missing Data and Practical Financial Examples

Type	Meaning	Finance Example
MCAR	Missing Completely at Random: Missingness is entirely independent of observed and unobserved variables.	Trade entry terminal intermittently experiences random network packet drops logging trades.
MAR	Missing at Random: Missingness is statistically related to other observed features in the dataset.	Younger retail banking customers are systematically more likely to leave the optional income field blank.
MNAR	Missing Not at Random: Missingness is directly correlated with the actual missing value itself.	Highly indebted distressed borrowers deliberately omit outstanding secondary credit obligations.

1.5.2 Imputation Methods

Common techniques for replacing missing data include: • Mean / Median / Mode, • K-Nearest Neighbors (KNN) Imputation, • Forward Fill, • Backward Fill, and • Interpolation.

Table 1.3: Advantages and Disadvantages of Imputation Methods

Method	Advantages	Disadvantages
Mean / Median	Simple, fast to compute	Distorts sample variance; ignores inter-feature relationships.
Mode	Directly applicable to categorical data	Can artificially over-represent the most common category.
KNN	Exploits multivariate relationships between features	Computationally expensive on large financial datasets.
Forward / Backward Fill	Well-suited for time-series sequences	Assumes most recent observed value remains representative.
Interpolation	Produces smooth estimates capturing trends	Unsuitable for non-sequential, unordered data.

1.5.3 Decision Framework for Selecting Imputation Methods

1.6 Imbalanced Categorical Data

Class imbalance arises when one category of the target variable appears far more frequently than other categories:




Figure 1.2: Decision Framework for Selecting Missing Value Imputation Methods.

Financial Example

In credit card fraud detection:

- Genuine Transactions: 99.9%
- Fraudulent Transactions: 0.1%

A naive model that blindly predicts “no fraud” every single time achieves 99.9% raw accuracy, yet is entirely useless. Thus, **accuracy alone is fundamentally misleading**.

1.6.1 Consequences of Class Imbalance

- Biases model parameters heavily toward majority class patterns.
- Renders overall accuracy metric uninformative for decision-makers.

1.6.2 Remedial Solutions

- **Oversampling:** Duplicating or synthetically generating minority class instances.
- **Undersampling:** Down-sampling majority class observations.
- **SMOTE:** Synthetic Minority Over-sampling Technique, interpolating new instances along minority feature lines.

1.7 Model Evaluation Metrics

For binary classification problems, predictions and actual outcomes are summarized in a confusion matrix:

Table 1.4: Binary Classification Confusion Matrix

	Predicted Positive	Predicted Negative
Actual Positive	True Positive (TP)	False Negative (FN)
Actual Negative	False Positive (FP)	True Negative (TN)

Precision

Of all instances predicted as positive, how many were genuinely correct?

$$\text{Precision} = \frac{TP}{TP + FP} \quad (1.3)$$

Recall

Of all actual positive instances in the population, how many did the model identify?

$$\text{Recall} = \frac{TP}{TP + FN} \quad (1.4)$$

 F_1 -Score

Harmonic mean of Precision and Recall, balancing false alarms against missed detections:

$$F_1 = 2 \cdot \frac{\text{Precision} \cdot \text{Recall}}{\text{Precision} + \text{Recall}} \quad (1.5)$$

1.8 Mistakes to be Avoided in ML Pipelines

1. **Global Scaling Before Splitting:** Scaling or normalizing the entire dataset prior to train/test partitioning causes data leakage.
2. **Blind Mean Imputation on MNAR Data:** Imputing missing values with unconditional means when data is Missing Not at Random injects substantial estimation bias.
3. **Ignoring Class Imbalance:** Relying exclusively on overall classification accuracy in skewed datasets conceals poor minority detection.
4. **Unchecked One-Hot Encoding on High-Cardinality Variables:** Creating dummy variables for high-cardinality categorical features leads to severe dimensionality inflation.
5. **Failing to Shuffle Non-Time-Series Data:** Neglecting random shuffling before partitioning standard cross-sectional data risks clustering-induced bias.

1.9 Key Takeaways

Core Principles of ML Data Preparation

- i) **Pre-processing dominates:** Data preprocessing accounts for 60% – 80% of total machine learning workflow effort.
- ii) **Always split first:** Partition data into training, validation, and test sets before performing any scaling or modeling.
- iii) **Diagnose missingness:** Determine whether missing values are MCAR, MAR, or MNAR before selecting an imputation strategy.
- iv) **Address class imbalance:** Use SMOTE, oversampling, or undersampling; accuracy alone is never sufficient.
- v) **Respect temporal order:** For financial and time-series data, strictly employ chrono-

logical splits and forward-looking imputation to prevent look-ahead bias.

Chapter 2

K-Nearest Neighbors (KNN)

2.1 Introduction to KNN

Definition 2.1.1: K-Nearest Neighbors

K-Nearest Neighbors (KNN) is a non-parametric, distance-based supervised machine learning algorithm used for both classification and regression tasks.

2.1.1 Core Algorithmic Characteristics

- **Supervised Learning:** Operates on labeled training datasets to predict continuous targets or discrete classes.
- **Dual Capability:** Utilized for classification (via majority voting) as well as regression (via local averaging).
- **Distance-Based Metric Method:** Assigns predictions purely based on geometric proximity in the feature space.
- **Lazy Learning / Instance-Based:** Does not construct an explicit internal model or parameter set during training; retains the entire training dataset and defers computation until inference time.

2.1.2 Algorithmic Workflow

The operational execution pipeline of KNN proceeds through four sequential stages:



Figure 2.1: Sequential Execution Pipeline of the k -Nearest Neighbors Algorithm.

2.2 Distance Metrics

Geometric closeness between query observations and stored instances is computed using formal metric distance functions.

2.2.1 Euclidean Distance

For two observations in a 2-dimensional feature space, $\mathbf{X} = (x_1, x_2)$ and $\mathbf{Y} = (y_1, y_2)$:

$$d(\mathbf{X}, \mathbf{Y}) = \sqrt{(x_1 - y_1)^2 + (x_2 - y_2)^2} \quad (2.1)$$

For an arbitrary p -dimensional feature vector (or n features):

$$d(\mathbf{X}, \mathbf{Y}) = \sqrt{\sum_{j=1}^p (x_j - y_j)^2} = \sqrt{\sum_{i=1}^n (x_i - y_i)^2} \quad (2.2)$$

2.2.2 Manhattan Distance

The Manhattan distance (L_1 norm) measures rectilinear city-block distance across coordinates:

$$d(\mathbf{X}, \mathbf{Y}) = \sum_{i=1}^n |x_i - y_i| \quad (2.3)$$

- **For 2 Dimensions:** $d = |x_1 - x_2| + |y_1 - y_2|$
- **For 3 Dimensions:** $d = |x_1 - x_2| + |y_1 - y_2| + |z_1 - z_2|$

2.2.3 Minkowski Distance

The generalized Minkowski metric of order $p \geq 1$ encompasses both Manhattan and Euclidean distances:

$$d(\mathbf{X}, \mathbf{Y}) = \left(\sum_{i=1}^n |x_i - y_i|^p \right)^{1/p} \quad (2.4)$$

Table 2.1: Minkowski Distance Orders and Resulting Metrics

Order p	Distance Metric	Mathematical Formula
$p = 1$	Manhattan distance	$\sum x_i - y_i $
$p = 2$	Euclidean distance	$\sqrt{\sum (x_i - y_i)^2}$
$p \geq 1$	Minkowski distance	$\left(\sum x_i - y_i ^p \right)^{1/p}$

Where:

- n = number of dimensions (features)
- p = positive real number (the “order” of the distance metric)
- x_i, y_i = the i -th feature coordinates of the two compared observations

2.3 Worked Numerical Examples

2.3.1 Fruit Classification via Euclidean Distance

Example 2.3.1: Fruit Classification (Apple vs. Lemon)

A dataset contains observations evaluated on two physical attributes: *Sweetness* and *Crunchiness*.

Fruit Training Dataset			
Point	Sweetness	Crunchiness	Class
<i>A</i>	7	7	Apple
<i>B</i>	8	6	Apple
<i>C</i>	3	8	Lemon
<i>D</i>	2	7	Lemon
<i>E</i>	6	8	Apple
<i>F</i>	1	9	Lemon

Task: Classify a new query observation $\mathbf{X} = (5, 7)$ using $k = 3$ and Euclidean distance.

Solution:

Step 1: Compute Euclidean distances from query point $\mathbf{X}(5, 7)$:

$$\text{To } A(7, 7) : d = \sqrt{(5 - 7)^2 + (7 - 7)^2} = \sqrt{4 + 0} = 2.00$$

$$\text{To } B(8, 6) : d = \sqrt{(5 - 8)^2 + (7 - 6)^2} = \sqrt{9 + 1} = \sqrt{10} \approx 3.16$$

$$\text{To } C(3, 8) : d = \sqrt{(5 - 3)^2 + (7 - 8)^2} = \sqrt{4 + 1} = \sqrt{5} \approx 2.23$$

$$\text{To } D(2, 7) : d = \sqrt{(5 - 2)^2 + (7 - 7)^2} = \sqrt{9 + 0} = 3.00$$

$$\text{To } E(6, 8) : d = \sqrt{(5 - 6)^2 + (7 - 8)^2} = \sqrt{1 + 1} = \sqrt{2} \approx 1.41$$

$$\text{To } F(1, 9) : d = \sqrt{(5 - 1)^2 + (7 - 9)^2} = \sqrt{16 + 4} = \sqrt{20} \approx 4.04$$

Step 2: Sort distances in ascending order:

Point	Distance	Class
<i>E</i>	1.41	Apple
<i>A</i>	2.00	Apple
<i>C</i>	2.23	Lemon
<i>D</i>	3.00	Lemon
<i>B</i>	3.16	Apple
<i>F</i>	4.04	Lemon

Step 3: Evaluate majority class for $k = 3$:

$$\{k = 3\} = \{E, A, C\}$$

- Apple: 2 votes (*E*, *A*)
- Lemon: 1 vote (*C*)

$\therefore \mathbf{X}$ is to be classified as **Apple**.

2.3.2 Fintech Stock Recommendation with Standardized Indicators

Example 2.3.2: Fintech Stock Recommendation (Buy, Hold, Avoid)

A fintech company builds a KNN system to recommend whether a stock is a **Buy**, **Hold**, or **Avoid** based on three financial indicators: Price-to-Earnings (P/E), Return on Equity (ROE), and Debt-to-Equity (D/E).

The training samples are provided after standardizing each indicator via $Z = \frac{x-\mu}{\sigma}$:

Standardized Financial Indicator Training Dataset

Stock	$P/E (Z)$	$ROE (Z)$	$D/E (Z)$	Class
S_1	-1.351	1.542	-1.184	Buy
S_2	-0.787	0.806	-0.858	Buy
S_3	-0.129	-0.299	-0.204	Hold
S_4	0.341	-0.668	0.286	Hold
S_5	1.750	-1.404	1.757	Avoid
S_6	1.093	-1.036	1.267	Avoid
S_7	-1.069	1.174	-1.021	Buy
S_8	0.153	-0.115	-0.041	Hold

Task: Predict the recommendation class for a new stock with standardized features:

$$\mathbf{X} = (-0.559, 0.437, -0.531)$$

using $k = 3$ and Euclidean distance.

Solution:

Step 1: Compute Euclidean distance for each stock:

$$d(\mathbf{X}, S_i) = \sqrt{(z_{1,X} - z_{1,S_i})^2 + (z_{2,X} - z_{2,S_i})^2 + (z_{3,X} - z_{3,S_i})^2}$$

Stock	$P/E (Z)$	$ROE (Z)$	$D/E (Z)$	Distance	Class
S_1	-1.351	1.542	-1.184	1.487	Buy
S_2	-0.787	0.806	-0.858	0.527	Buy
S_3	-0.129	-0.299	-0.204	0.939	Hold
S_4	0.341	-0.668	0.286	1.266	Hold
S_5	1.750	-1.404	1.757	3.760	Avoid
S_6	1.093	-1.036	1.267	2.874	Avoid
S_7	-1.069	1.174	-1.021	1.002	Buy
S_8	0.153	-0.115	-0.041	1.053	Hold

Step 2: Sorted Euclidean distance table:

Stock	Distance	Class
S_2	0.527	Buy
S_3	0.939	Hold
S_7	1.002	Buy
S_8	1.053	Hold
S_4	1.266	Hold
S_1	1.487	Buy
S_6	2.874	Avoid
S_5	3.760	Avoid

Step 3: Majority voting for $k = 3$:

$$\{k = 3\} = \{S_2, S_3, S_7\}$$

- Buy: 2 votes (S_2, S_7)
- Hold: 1 vote (S_3)
- Avoid: 0 votes

$\therefore$ The new point is to be classified as **Buy** for $k = 3$.

2.3.3 Bank Loan Approval via Manhattan Distance

Example 2.3.3: Bank Loan Approval Decision

A bank utilizes KNN to decide whether to **Approve** or **Reject** a loan application based on two features:

- Credit score (scaled by dividing by 100, such that $750 \implies 7.5$)
- Annual Income (in Rs. Lakhs)

Loan Applicant Training Dataset

Applicant	Credit Score	Annual Income (Rs. Lakhs)	Decision
A_1	7.5	8.0	Approve
A_2	6.8	6.0	Approve
A_3	6.0	4.0	Reject
A_4	5.5	3.0	Reject
A_5	7.2	7.0	Approve
A_6	5.8	3.5	Reject

Task: Predict the decision for a new applicant with Credit Score = 6.5 and Annual Income = 5.5 Lakhs, i.e., $\mathbf{X} = (6.5, 5.5)$, using $k = 3$ and Manhattan distance.

Solution:

Step 1: Compute Manhattan distances from query point $\mathbf{X}(6.5, 5.5)$:

$$\text{To } A_1(7.5, 8.0) : d = |6.5 - 7.5| + |5.5 - 8.0| = 1.0 + 2.5 = 3.5$$

$$\text{To } A_2(6.8, 6.0) : d = |6.5 - 6.8| + |5.5 - 6.0| = 0.3 + 0.5 = 0.8$$

$$\text{To } A_3(6.0, 4.0) : d = |6.5 - 6.0| + |5.5 - 4.0| = 0.5 + 1.5 = 2.0$$

$$\text{To } A_4(5.5, 3.0) : d = |6.5 - 5.5| + |5.5 - 3.0| = 1.0 + 2.5 = 3.5$$

$$\text{To } A_5(7.2, 7.0) : d = |6.5 - 7.2| + |5.5 - 7.0| = 0.7 + 1.5 = 2.2$$

$$\text{To } A_6(5.8, 3.5) : d = |6.5 - 5.8| + |5.5 - 3.5| = 0.7 + 2.0 = 2.7$$

Step 2: Sorted Manhattan distance ranking:

Applicant	Distance	Decision
A_2	0.8	Approve
A_3	2.0	Reject
A_5	2.2	Approve
A_6	2.7	Reject
A_4	3.5	Reject
A_1	3.5	Approve

Step 3: Majority vote for $k = 3$:

$$\{k = 3\} = \{A_2, A_3, A_5\}$$

- Approve: 2 votes (A_2, A_5)
- Reject: 1 vote (A_3)

$\therefore$ The new applicant $\mathbf{X} = (6.5, 5.5)$ is to be concluded as **Approve** when $k = 3$.

2.3.4 Real Estate House Price Prediction (KNN Regression)**Example 2.3.4: Real Estate Price Estimation (KNN Regression)**

A real estate agency uses KNN regression to estimate house prices based on:

- Area (in hundreds of sq. ft., e.g., 1400 sq. ft. $\implies$ 14)
- Age (in years)

Real Estate Housing Dataset

House	Area (hundred sq. ft.)	Age (years)	Price (Rs. Lakhs)
H_1	12	5	50
H_2	15	3	65
H_3	8	10	30
H_4	20	2	90
H_5	10	8	38
H_6	18	6	75

Task: Predict the price of a new house with Area = 14 (1400 sq. ft.) and Age = 4 years, i.e., $\mathbf{X} = (14, 4)$, using $k = 3$ and Manhattan distance.

Solution:

Step 1: Compute Manhattan distances from query house $\mathbf{X}(14, 4)$:

$$\begin{aligned} \text{To } H_1(12, 5) : d &= |14 - 12| + |4 - 5| = 2 + 1 = 3 \\ \text{To } H_2(15, 3) : d &= |14 - 15| + |4 - 3| = 1 + 1 = 2 \\ \text{To } H_3(8, 10) : d &= |14 - 8| + |4 - 10| = 6 + 6 = 12 \\ \text{To } H_4(20, 2) : d &= |14 - 20| + |4 - 2| = 6 + 2 = 8 \\ \text{To } H_5(10, 8) : d &= |14 - 10| + |4 - 8| = 4 + 4 = 8 \\ \text{To } H_6(18, 6) : d &= |14 - 18| + |4 - 6| = 4 + 2 = 6 \end{aligned}$$

Step 2: Sorted Manhattan distance ranking:

House	Distance	Price (Rs. Lakhs)
H_2	2	65
H_1	3	50
H_6	6	75
H_4	8	90
H_5	8	38
H_3	12	30

Step 3: Average prices across $k = 3$ nearest neighbors:

$$\begin{aligned} \{k = 3\} &= \{H_2, H_1, H_6\} \\ \hat{y} &= \frac{65 + 50 + 75}{3} = \frac{190}{3} \approx 63.33 \text{ Lakhs} \end{aligned}$$

$\therefore$ The new point $\mathbf{X} = (14, 4)$ is predicted to have a price of **Rs. 63.33 Lakhs** for $k = 3$.

2.4 Parameter Selection: Small k vs. Large k

Selecting the hyperparameter k balances local sensitivity against global regularization:

Table 2.2: Comparison of Small k versus Large k Characteristics

Aspect	Small k	Large k
Neighbourhood	Local	Broad
Sensitivity to Noise	High	Lower
Variance	Higher	Lower
Bias	Lower	Higher
Decision Boundary	Irregular / complex	Smoother
Risk	Overfitting	Underfitting

Guiding Principle for Choosing k

Choose k to balance bias and variance.

Chapter 3

Naive Bayes Classification

3.1 Probability Foundations & Notations

The Naive Bayes classifier is grounded in classical axiomatic probability theory and conditional distributions.

Table 3.1: Probability Notations and Terminology

Notation	Meaning
$P(A)$	Probability of event A
$P(A^c)$	Probability of complement of A
$P(A \cap B)$	Joint probability of A and B occurring simultaneously
$P(A \cup B)$	Probability of either event A or event B occurring
$P(A B)$	Conditional probability of event A given event B has occurred
$P(A, B)$	Joint probability of events A and B (alternative standard notation)

3.1.1 Key Probability Identities

- **Complement Rule:**

$$P(A^c) = 1 - P(A) \quad (3.1)$$

- **Addition Rule (Union):**

$$P(A \cup B) = P(A) + P(B) - P(A \cap B) \quad (3.2)$$

- **Joint Probability Equivalence:**

$$P(A \cap B) = P(A, B) \quad (3.3)$$

3.1.2 Conditional Probability & Multiplication Rule

Definition 3.1.1: Conditional Probability

For two events A and B with $P(B) > 0$:

$$P(A | B) = \frac{P(A \cap B)}{P(B)} \quad (3.4)$$

Re-arranging yields the general **Multiplication Rule**:

$$P(A, B) = P(A | B)P(B) \quad (3.5)$$

$$P(A, B) = P(B | A)P(A) \quad (3.6)$$

3.1.3 Statistical Independence

If events A and B are statistically independent:

$$P(A \cap B) = P(A) \cdot P(B) \quad (3.7)$$

and equivalently:

$$P(A | B) = P(A) \quad (3.8)$$

3.2 Bayes' Theorem

Suppose an analyst wishes to determine the probability that an observation belongs to class C_k after observing feature evidence vector $\mathbf{X}$:

Theorem 3.2.1: Bayes' Theorem

$$P(C_k | \mathbf{X}) = \frac{P(\mathbf{X} | C_k)P(C_k)}{P(\mathbf{X})} \quad (3.9)$$

In conceptual terms:

$$\text{Posterior} = \frac{\text{Likelihood} \times \text{Prior}}{\text{Evidence}} \quad (3.10)$$

3.3 Worked Numerical Example: Credit Risk Classification

Example 3.3.1: Credit Risk Classification

A retail lending institution assesses credit risk with two target states:

- C_1 : Good Credit
- C_2 : Default

Historical portfolio priors and likelihoods for a financial indicator vector $\mathbf{X}$ are given by:

$$\begin{aligned} P(C_1) &= 0.80, & P(C_2) &= 0.20 \\ P(\mathbf{X} | C_1) &= 0.10, & P(\mathbf{X} | C_2) &= 0.60 \end{aligned}$$

Task: Compute the evidence $P(\mathbf{X})$, determine the posterior probabilities $P(C_1 | \mathbf{X})$ and $P(C_2 | \mathbf{X})$, and classify the borrower.

Solution:

Step 1: Compute Total Evidence $P(\mathbf{X})$ via Law of Total Probability:

$$\begin{aligned} P(\mathbf{X}) &= P(\mathbf{X} | C_1)P(C_1) + P(\mathbf{X} | C_2)P(C_2) \\ &= (0.10 \times 0.80) + (0.60 \times 0.20) \\ &= 0.08 + 0.12 \\ &= 0.20 \end{aligned}$$

Step 2: Calculate Posterior Probabilities:

$$P(C_1 | \mathbf{X}) = \frac{P(\mathbf{X} | C_1)P(C_1)}{P(\mathbf{X})} = \frac{0.10 \times 0.80}{0.20} = \frac{0.08}{0.20} = 0.40$$

$$P(C_2 | \mathbf{X}) = \frac{P(\mathbf{X} | C_2)P(C_2)}{P(\mathbf{X})} = \frac{0.60 \times 0.20}{0.20} = \frac{0.12}{0.20} = 0.60$$

Step 3: Comparison and Decision:

$$P(C_2 | \mathbf{X}) > P(C_1 | \mathbf{X}) \quad (0.60 > 0.40)$$

$\therefore$ Classify $\mathbf{X}$ as **Default** (C_2).

3.4 The Bayes Classifier Framework

3.4.1 Marginal Evidence Across K Classes

For a classification problem with K discrete classes $C_1, C_2, \dots, C_K$, the overall marginal probability of evidence $\mathbf{X}$ is:

$$P(\mathbf{X}) = \sum_{k=1}^K P(\mathbf{X} | C_k)P(C_k) \quad (3.11)$$

3.4.2 Optimal Bayes Decision Rule

The Bayes classifier assigns an observation $\mathbf{X}$ to the class maximizing the posterior probability:

$$\hat{c} = \arg \max_{C_k} P(C_k | \mathbf{X}) = \arg \max_{C_k} \frac{P(\mathbf{X} | C_k)P(C_k)}{P(\mathbf{X})} \quad (3.12)$$

Because the evidence denominator $P(\mathbf{X})$ is identical across all classes C_k , the decision rule simplifies to:

$$\hat{c} = \arg \max_{C_k} P(\mathbf{X} | C_k)P(C_k) \quad (3.13)$$

3.5 The Naive Bayes Classifier

3.5.1 The Challenge of High-Dimensional Likelihoods

When an observation is characterized by p features $\mathbf{X} = (x_1, x_2, \dots, x_p)$:

$$P(C_k | x_1, x_2, \dots, x_p) = \frac{P(x_1, x_2, \dots, x_p | C_k)P(C_k)}{P(x_1, x_2, \dots, x_p)} \quad (3.14)$$

Directly estimating the joint multivariate conditional distribution $P(x_1, x_2, \dots, x_p | C_k)$ from finite sample sizes is computationally and statistically challenging due to exponential data sparsity.

3.5.2 The Naive Conditional Independence Assumption

Definition 3.5.1: The Naive Assumption

All input features $x_1, x_2, \dots, x_p$ are mutually conditionally independent given the class label C_k :

$$P(x_1, x_2, \dots, x_p \mid C_k) = \prod_{j=1}^p P(x_j \mid C_k) \quad (3.15)$$

that is:

$$P(x_1, x_2, \dots, x_p \mid C_k) = P(x_1 \mid C_k) \cdot P(x_2 \mid C_k) \cdots P(x_p \mid C_k) \quad (3.16)$$

3.5.3 Naive Bayes Classification Rule

Applying the independence assumption produces the core operational classification formula:

$$\hat{c} = \arg \max_{C_k} P(C_k) \prod_{j=1}^p P(x_j \mid C_k) \quad (3.17)$$

Why Is It Termed “Naive”?

The assumption is termed “naive” because conditional independence among predictors is rarely strictly true in real-world scenarios.

Financial Example: In credit scoring, variables such as *Income*, *Debt*, and *Employment Status* clearly exhibit mutual dependencies. However, in practice, the naive approximation dramatically simplifies computation while delivering strong empirical classification performance.

3.5.4 Algorithmic Step-by-Step Sequence

Naive Bayes Execution Procedure

- i) **Start with Training Data:** Assemble feature vectors $\mathbf{X}_i$ and class labels y_i .
- ii) **Calculate Prior Probabilities:** Compute class base rates $P(C_k) = \frac{N_{C_k}}{N}$.
- iii) **Calculate Univariate Likelihoods:** Estimate 1-D conditional probabilities $P(x_j | C_k)$ for each feature given each class.
- iv) **Apply Naive Assumption:** Decompose joint likelihood into the product:

$$P(x_1, \dots, x_p | C_k) = \prod_{j=1}^p P(x_j | C_k)$$

- v) **Calculate Class Score:** Evaluate the unnormalized joint score for each candidate class:

$$\text{Score}(C_k) = P(C_k) \prod_{j=1}^p P(x_j | C_k)$$

- vi) **Compare Scores Across All Classes:** Evaluate $\text{Score}(C_k)$ across $k = 1, \dots, K$.
- vii) **Predict Output Class:**

$$\hat{c} = \arg \max_{C_k} \text{Score}(C_k)$$

Chapter 4

Logistic Regression

4.1 Introduction and Core Motivation

Logistic regression is a fundamental supervised learning algorithm utilized for classification problems. Unlike regression models that predict a continuous numerical quantity, logistic regression models the conditional probability that a given observation belongs to a particular class.

4.1.1 Why Not Standard Linear Regression?

In ordinary linear regression, the relationship between features and the target is modeled as:

$$Y = \beta_0 + \beta_1 X_1 + \cdots + \beta_p X_p + \varepsilon \quad (4.1)$$

Under this linear formulation, the predicted value $\hat{Y}$ can take any real value across the entire unbounded interval:

$$\hat{Y} \in (-\infty, +\infty) \quad (4.2)$$

However, probabilities are mathematically constrained to lie strictly within the unit interval:

$$P(Y = 1 \mid X) \in [0, 1] \quad (4.3)$$

If ordinary linear regression were used for classification, predictions could yield nonsensical probabilities less than 0 or greater than 1. Consequently, a transformation is needed that “squashes” the unbounded linear output $(-\infty, +\infty)$ into the valid probability range $[0, 1]$. In logistic regression, this transformation is provided by the **sigmoid (logistic) function**.

4.2 The Sigmoid (Logistic) Function and Probability Modeling

Logistic regression first computes a linear combination of the input features:

$$Z = \beta_0 + \beta_1 X_1 + \cdots + \beta_p X_p \quad (4.4)$$

The real-valued score Z is then passed through the sigmoid (logistic) activation function to yield the predicted probability of the positive class ($Y = 1$):

$$P(Y = 1) = \frac{1}{1 + e^{-Z}} \quad (4.5)$$

4.2.1 Algebraic Equivalence of the Sigmoid Function

Multiplying both the numerator and the denominator by e^Z (since $\frac{e^Z}{e^Z} = 1$) provides the alternate standard representation:

$$P(Y = 1) = \frac{1 \cdot e^Z}{(1 + e^{-Z}) \cdot e^Z} \quad (4.6)$$

$$= \frac{e^Z}{e^Z + e^{-Z} \cdot e^Z} \quad (4.7)$$

$$= \frac{e^Z}{e^Z + e^{Z-Z}} \quad (4.8)$$

$$= \frac{e^Z}{e^Z + e^0} \quad (4.9)$$

$$= \frac{e^Z}{e^Z + 1} = \frac{e^Z}{1 + e^Z} \quad (4.10)$$

Substituting the linear combination $Z = \beta_0 + \beta_1 X_1 + \dots + \beta_p X_p$ directly yields:

$$P(Y = 1) = \frac{e^{\beta_0 + \beta_1 X_1 + \dots + \beta_p X_p}}{1 + e^{\beta_0 + \beta_1 X_1 + \dots + \beta_p X_p}} \quad (4.11)$$

4.2.2 Properties of the Sigmoid Curve and Decision Threshold

The sigmoid curve possesses distinctive geometric characteristics:

- **S-Shaped Geometry:** The curve exhibits a smooth, monotonic S-shape.
- **Asymptotic Bounds:** As $Z \rightarrow -\infty$, $P(Y = 1) \rightarrow 0$. As $Z \rightarrow +\infty$, $P(Y = 1) \rightarrow 1$.
- **Symmetric Midpoint:** At $Z = 0$, $e^0 = 1$, yielding exactly $P(Y = 1) = \frac{1}{1+1} = 0.5$.

To convert continuous probability estimates into discrete class labels, a **decision threshold** is established (conventionally set to 0.5):

$$\hat{Y} = \begin{cases} 1, & \text{if } P(Y = 1) \geq 0.5 \quad (Z \geq 0) \\ 0, & \text{if } P(Y = 1) < 0.5 \quad (Z < 0) \end{cases} \quad (4.12)$$




Figure 4.1: The Sigmoid activation function squashing unbounded linear values Z into valid class probabilities $[0, 1]$ with a decision boundary at $Z = 0$ ($P = 0.5$).

4.3 Key Terms and Matrix Notation

Definition 4.3.1: Core Variables of Logistic Regression

- **Independent Variables (X):** The set of input features or predictor variables $X_1, X_2, \dots, X_p$.
- **Dependent Variable (Y):** The categorical target outcome, typically binary $Y \in \{0, 1\}$.
- **Logistic Function (Sigmoid Function):** The nonlinear squashing function $\sigma(Z) = \frac{1}{1+e^{-Z}}$ mapping real scores to $(0, 1)$.
- **Coefficient (β_j):** The parameter estimated by the model for predictor X_j , indicating how strongly X_j affects the dependent variable on the log-odds scale.
- **Intercept (β_0):** The constant baseline term representing the log-odds when all independent variables equal zero ($X_1 = X_2 = \dots = X_p = 0$).

4.3.1 Matrix Representation

In compact linear algebra notation, the dataset of n observations and p features is represented as an $n \times p$ feature matrix X :

$$X = \begin{bmatrix} x_{11} & x_{12} & \dots & x_{1p} \\ x_{21} & x_{22} & \dots & x_{2p} \\ \vdots & \vdots & \ddots & \vdots \\ x_{n1} & x_{n2} & \dots & x_{np} \end{bmatrix} \quad (4.13)$$

where:

- n is the total number of observations (samples).
- p is the total number of features (predictors).

The observed class labels form a binary response vector $y \in \{0, 1\}^n$:

$$y = \begin{bmatrix} y_1 \\ y_2 \\ \vdots \\ y_n \end{bmatrix}, \quad y_i \in \{0, 1\} \quad (4.14)$$

4.4 Odds, Log-Odds (Logit), and Parameter Interpretation

4.4.1 Definition of Odds

The **odds** of an event represent the ratio of the probability that the event occurs to the probability that the event does not occur:

$$\text{Odds} = \frac{P(Y = 1)}{1 - P(Y = 1)} \quad (4.15)$$

Example 4.4.1: Calculating Probabilities and Odds

Consider a basket containing 3 apples and 5 oranges (total items = 3 + 5 = 8).

1. Probability of selecting an Orange:

$$P(\text{Orange}) = \frac{5}{8} = 0.625 \quad (4.16)$$

2. Odds of selecting an Orange:

$$\text{Odds}(\text{Orange}) = \frac{\text{Number of Oranges}}{\text{Number of Not Oranges (Apples)}} = \frac{5}{3} \approx 1.667 \quad (4.17)$$

3. Odds of selecting an Apple:

$$\text{Odds}(\text{Apple}) = \frac{\text{Number of Apples}}{\text{Number of Not Apples (Oranges)}} = \frac{3}{5} = 0.600 \quad (4.18)$$

4.4.2 Log-Odds (Logit Function)

Taking the natural logarithm of the odds produces the **log-odds** (or **logit**):

$$\ln(\text{Odds}) = \ln \left(\frac{P(Y = 1)}{1 - P(Y = 1)} \right) \quad (4.19)$$

Substituting $P(Y = 1) = \frac{e^Z}{1+e^Z}$ and $1 - P(Y = 1) = \frac{1}{1+e^Z}$:

$$\frac{P(Y = 1)}{1 - P(Y = 1)} = \frac{\frac{e^Z}{1+e^Z}}{\frac{1}{1+e^Z}} = e^Z \quad (4.20)$$

Taking the natural logarithm:

$$\ln \left(\frac{P(Y = 1)}{1 - P(Y = 1)} \right) = \ln(e^Z) = Z = \beta_0 + \beta_1 X_1 + \cdots + \beta_p X_p \quad (4.21)$$

Thus, **logistic regression is a linear model for the log-odds**.

4.4.3 Interpreting Coefficients and Odds Ratios

For a one-unit increase in predictor X_j (holding all other variables constant), the log-odds changes by β_j :

$$\Delta \ln(\text{Odds}) = \beta_j \quad (4.22)$$

Exponentiating both sides yields the **Odds Ratio (OR)**:

$$\text{Odds Ratio} = e^{\beta_j} \quad (4.23)$$

The directional impact of β_j on the odds is summarized as follows:

- $\beta_j > 0 \implies e^{\beta_j} > 1$: The odds of the event increase.
- $\beta_j < 0 \implies e^{\beta_j} < 1$: The odds of the event decrease.
- $\beta_j = 0 \implies e^{\beta_j} = 1$: Predictor X_j has no effect on the odds.

Example 4.4.2: Odds Ratio Interpretation with $\beta_j = 0.7$

Suppose the estimated coefficient for a predictor X_j is $\beta_j = 0.7$.

$$\text{Odds Ratio} = e^{\beta_j} = e^{0.7} \approx 2.0137 \approx 2 \quad (4.24)$$

Interpretation: A one-unit increase in X_j roughly doubles the odds of the outcome event occurring ($Y = 1$).

4.5 Model Training via Maximum Likelihood Estimation (MLE)

Logistic regression estimates its parameters $\beta_0, \beta_1, \dots, \beta_p$ using **Maximum Likelihood Estimation (MLE)**.

- **Objective of MLE:** Find the parameter vector β that maximizes the joint likelihood of observing the given sample data.
- **Lack of Closed-Form Solution:** Unlike ordinary least squares (OLS) regression where $\hat{\beta} = (X^T X)^{-1} X^T y$, there is **no closed-form analytical solution** for the MLE of logistic regression parameters.
- **Numerical Optimization:** Parameter estimates are solved numerically using iterative optimization algorithms (e.g., Newton-Raphson, Iteratively Reweighted Least Squares, or Gradient Ascent).

4.6 Assumptions of Logistic Regression

Valid inference and stable performance in logistic regression rely on six foundational assumptions:

1. **Independent Observations:** Sample observations must be independent of one another (no repeated measures or auto-correlation).
2. **Binary Dependent Variable:** The response variable Y is binary (for standard binomial logistic regression).
3. **Linearity in Log-Odds:** Independent variables exhibit a strictly linear relationship with the log-odds $\ln\left(\frac{P}{1-P}\right)$.

4. **Absence of Extreme Outliers:** The predictor space should not contain extreme leverage points or anomalous outliers.
5. **Adequately Large Sample Size:** Because MLE relies on large-sample asymptotic properties, sufficient sample size is required to prevent biased parameter estimates.
6. **Low Multicollinearity:** Predictor variables should exhibit little to no multicollinearity. High correlation among predictors inflates parameter standard errors.

4.7 Variants, Metrics, and Trade-Offs

4.7.1 Types of Logistic Regression

Depending on the nature of the target variable, logistic regression is categorized into three types:

- **Binomial Logistic Regression:** The target has exactly two classes (e.g., Yes / No, Default / Non-Default).
- **Multinomial Logistic Regression:** The target has three or more unordered nominal classes (e.g., Cat, Dog, Sheep).
- **Ordinal Logistic Regression:** The target has three or more ordered categories (e.g., Low, Medium, High).

4.7.2 Model Evaluation Metrics

Model performance for logistic regression is assessed using standard classification metrics:

- **Accuracy:** Ratio of correctly classified instances to total instances.
- **Precision:** Ratio of true positives to all predicted positives ($\frac{TP}{TP+FP}$).
- **Recall (Sensitivity):** Ratio of true positives to all actual positives ($\frac{TP}{TP+FN}$).
- **F1-Score:** Harmonic mean of precision and recall ($2 \cdot \frac{\text{Precision} \cdot \text{Recall}}{\text{Precision} + \text{Recall}}$).
- **AUC-ROC:** Area Under the Receiver Operating Characteristic curve, measuring discrimination across all possible classification thresholds.

4.7.3 Advantages and Disadvantages

Advantages (Pros)	Disadvantages (Cons)
• Outputs well-calibrated probabilities, not just hard class labels. • Highly interpretable parameters via odds ratios (e^{β_j}). • Directly extends to multiclass classification problems.	• Interaction terms between features must be constructed manually. • Interpretations are multiplicative on odds, not linear on probabilities. • Complete separation prevents convergence (coefficients diverge to $\pm\infty$). • Requires regularization (e.g., L1/L2 penalties) or informative priors to handle separation.

4.8 End-to-End Workflow and Generalized Linear Model (GLM) Classification

4.8.1 Five-Step Implementation Pipeline

1. **Prepare Data:** Encode categorical predictors (e.g., one-hot or dummy encoding) and scale continuous features if regularization is utilized.
2. **Fit Model:** Estimate parameter vector β using Maximum Likelihood Estimation (MLE) via numerical optimization.
3. **Interpret Coefficients:** Exponentiate parameters to interpret predictor effects as odds ratios (e^{β_j}).
4. **Predict and Threshold:** Compute predicted probabilities $P(Y = 1) = \sigma(X\beta)$ and assign class labels based on a selected threshold (typically 0.5).
5. **Evaluate Model:** Measure performance using Accuracy, Precision, Recall, F1-Score, and AUC-ROC.

4.8.2 Membership in Generalized Linear Models (GLM)

Logistic regression is a member of the broader **Generalized Linear Models (GLM)** family. In GLM theory, a model is defined by three components:

1. **Random Component:** The dependent variable Y follows a Bernoulli (or Binomial) distribution: $Y \sim \text{Bernoulli}(p)$.
2. **Systematic Component:** A linear combination of predictors $Z = \beta_0 + \sum_{j=1}^p \beta_j X_j$.
3. **Link Function:** The **logit link function** $g(p) = \ln\left(\frac{p}{1-p}\right) = Z$, which maps the expected probability parameter from $[0, 1]$ to the unbounded real line $(-\infty, +\infty)$.

Chapter 5

Decision Trees

5.1 Overview and Core Concepts

A **Decision Tree** is a supervised machine learning model that predicts a target outcome by recursively splitting the data into smaller, progressively more homogeneous (pure) subsets based on feature values.

5.1.1 Structural Components of a Decision Tree

- **Root Node:** The topmost node of the tree, containing all training samples prior to any partitioning.
- **Decision Node (Internal Node):** A node that tests a specific feature condition to split data into child branches.
- **Branch:** A directed link representing the outcome of a decision test leading to a subsequent node.
- **Leaf Node (Terminal Node):** An end node carrying the final prediction (class label in classification or mean value in regression); it undergoes no further splitting.
- **Parent Node:** A node immediately preceding a split.
- **Child Node:** A node generated immediately as the outcome of a parent node split.
- **Tree Depth:** The maximum number of levels (edges) from the root node to the deepest leaf node.

5.1.2 General Recursive Splitting Algorithm

1. **Initialization:** Place all training observations at the root node.
2. **Check Termination:** If a stopping condition is met (e.g., node is pure, maximum depth reached, or insufficient samples), convert the node into a terminal leaf.
3. **Evaluate Candidate Splits:** For every available feature and potential split point, compute the impurity reduction across the subsets.
4. **Select Optimal Split:** Select the feature and threshold that yields the best split criterion:
 - Maximum Information Gain (ID3 algorithm),
 - Maximum Gain Ratio (C4.5 algorithm),
 - Maximum Gini Gain / Minimum Weighted Gini Impurity (CART algorithm), or

- Maximum Variance Reduction (for continuous regression trees).
5. **Recursive Partitioning:** Partition the sample observations into disjoint child subsets and recurse on each child node.
6. **Stopping Criteria:** Terminate tree growth when:
- The node is completely pure (all samples share the same class),
 - The predefined maximum depth has been reached,
 - The minimum sample size required for a split is breached, or
 - No available candidate split yields a positive impurity improvement.



Figure 5.1: End-to-end split evaluation pipeline for constructing decision tree nodes.

5.2 Impurity Measures and Splitting Criteria

5.2.1 Standard Tree Induction Algorithms

Different decision tree induction algorithms utilize distinct impurity metrics and optimization objectives:

Algorithm	Impurity Measure	Split Criterion
ID3	Entropy	Information Gain
C4.5	Entropy	Gain Ratio
CART	Gini Impurity	Gini Gain / Minimum Weighted Gini

5.2.2 Entropy and Information Gain

Entropy quantifies the amount of uncertainty or disorder in a dataset S containing c classes:

$$H(S) = - \sum_{i=1}^c p_i \log_2(p_i) \quad (5.1)$$

where $p_i = \frac{\text{Number of samples in class } i}{\text{Total observations in node}}$ is the empirical class probability.

For a two-class problem ($c = 2$ with probabilities p_1 and p_2):

$$H(S) = -p_1 \log_2(p_1) - p_2 \log_2(p_2) \quad (5.2)$$

- A node containing only one class is completely pure: $H(S) = 0$.
- Maximum entropy occurs under equal class balance ($p_1 = p_2 = 0.5 \implies H(S) = 1.0$).

Information Gain (IG) measures the reduction in entropy achieved by partitioning a parent dataset S using attribute A :

$$IG(S, A) = H(\text{parent}) - H(\text{children}) = H(S) - \sum_{v \in \text{Values}(A)} \frac{|S_v|}{|S|} H(S_v) \quad (5.3)$$

where $H(\text{children}) = \sum_i \frac{n_i}{n} H_i$ represents the weighted average entropy across all resulting child partitions S_v of size n_i .

5.2.3 Gain Ratio (C4.5)

To overcome ID3's bias toward attributes with many distinct values, the **C4.5** algorithm normalizes Information Gain by the split information:

$$\text{Gain Ratio}(S, A) = \frac{IG(S, A)}{\text{Split Info}(S, A)} \quad (5.4)$$

where $\text{Split Info}(S, A) = -\sum_{v \in \text{Values}(A)} \frac{|S_v|}{|S|} \log_2 \left(\frac{|S_v|}{|S|} \right)$.

5.2.4 Gini Impurity (CART)

Gini Impurity measures the probability of misclassifying a randomly chosen element from the set if it were randomly labeled according to the class distribution:

$$\text{Gini}(S) = 1 - \sum_{i=1}^c p_i^2 \quad (5.5)$$

For a binary classification setting ($c = 2$):

$$\text{Gini}(S) = 1 - p_1^2 - p_2^2 \quad (5.6)$$

- **Smaller Gini:** Indicates a purer node.
- **Larger Gini:** Indicates a more mixed, impure node.
- $\text{Gini} = 0 \implies$ The node is completely pure (all samples belong to a single class).
- For binary classification, maximum Gini is 0.5 when $p_1 = p_2 = 0.5$.

5.2.5 Comparison: Gini Impurity vs. Entropy

Property	Gini Impurity	Entropy
Measures	Impurity	Uncertainty
Minimum Value	0	0
Pure Node Value	0	0
Binary Maximum	0.5 (at $p = 0.5$)	1.0 (at $p = 0.5$)
Good Split Condition	Lower Weighted Gini	Higher Information Gain
Computation	Simpler (algebraic squares)	Heavy (evaluates logarithms)
Common Usage	CART	ID3 / C4.5 algorithms

5.2.6 Advantages and Disadvantages of Decision Trees

Advantages	Disadvantages
• Highly intuitive and easy to understand. • Easy to visualize graphically. • Easy to explain to non-technical audiences. • Captures complex non-linear relationships. • Directly applicable to classification and regression.	• Highly susceptible to overfitting on training data. • Deep trees become unwieldy and overly complex. • High variance: small data perturbations alter the tree. • A single tree has limited predictive stability.

5.3 Solved Problem 5.1: Play Tennis Decision Tree

Example 5.3.1: ID3 Induction on Play Tennis Dataset

Construct an optimal decision tree using the ID3 algorithm (Entropy and Information Gain) on the classic 15-sample Play Tennis dataset.

Day	Outlook	Temp	Humidity	Wind	Tennis?
1	Sunny (S)	Hot (H)	High (H)	Weak (W)	No (N)
2	Sunny (S)	Hot (H)	High (H)	Strong (S)	No (N)
3	Overcast (O)	Hot (H)	High (H)	Weak (W)	Yes (Y)
4	Rain (R)	Mild (M)	High (H)	Weak (W)	Yes (Y)
5	Rain (R)	Cool (C)	Normal (N)	Weak (W)	Yes (Y)
6	Rain (R)	Cool (C)	Normal (N)	Strong (S)	No (N)
7	Overcast (O)	Cool (C)	Normal (N)	Strong (S)	Yes (Y)
8	Sunny (S)	Mild (M)	High (H)	Weak (W)	No (N)
9	Sunny (S)	Cool (C)	Normal (N)	Weak (W)	Yes (Y)
10	Rain (R)	Mild (M)	Normal (N)	Weak (W)	Yes (Y)
11	Sunny (S)	Mild (M)	Normal (N)	Strong (S)	Yes (Y)
12	Overcast (O)	Mild (M)	High (H)	Strong (S)	Yes (Y)
13	Overcast (O)	Hot (H)	Normal (N)	Weak (W)	Yes (Y)
14	Rain (R)	Mild (M)	High (H)	Strong (S)	No (N)
15	Sunny (S)	Hot (H)	High (H)	Weak (W)	No (N)

Overall Dataset Statistics and Entropy

The target variable contains 15 total observations: 9 Yes and 6 No:

$$p_1 = P(Y) = \frac{9}{15} = 0.60, \quad p_2 = P(N) = \frac{6}{15} = 0.40 \quad (5.7)$$

$$H(S) = -0.6 \log_2(0.6) - 0.4 \log_2(0.4) = 0.97095 \approx 0.9709 \quad (5.8)$$

Root Attribute Evaluation

We calculate the weighted entropy $H(\text{children}) = \sum \frac{n_i}{n} H_i$ for each candidate predictor:

1. Outlook ($n = 15$):

Value	Count (n_i)	Yes (Y)	No (N)	Entropy (H_i)	Weighted Entropy
Sunny (S)	6	2	4	0.9182	$\frac{6}{15} \times 0.9182 = 0.3673$
Overcast (O)	4	4	0	0.0000	$\frac{4}{15} \times 0 = 0.0000$
Rain (R)	5	3	2	0.9709	$\frac{5}{15} \times 0.9709 = 0.3236$
Total	15	9	6		0.6909

2. Temperature ($n = 15$):

Value	Count (n_i)	Yes (Y)	No (N)	Entropy (H_i)	Weighted Entropy
Hot (H)	5	2	3	0.9709	$\frac{5}{15} \times 0.9709 = 0.3236$
Mild (M)	6	4	2	0.9182	$\frac{6}{15} \times 0.9182 = 0.3673$
Cool (C)	4	3	1	0.8112	$\frac{4}{15} \times 0.8112 = 0.2163$
Total	15	9	6		0.9072

3. Humidity ($n = 15$):

Value	Count (n_i)	Yes (Y)	No (N)	Entropy (H_i)	Weighted Entropy
High (H)	8	3	5	0.9544	$\frac{8}{15} \times 0.9544 = 0.5090$
Normal (N)	7	6	1	0.5916	$\frac{7}{15} \times 0.5916 = 0.2761$
Total	15	9	6		0.7851

4. Wind ($n = 15$):

Value	Count (n_i)	Yes (Y)	No (N)	Entropy (H_i)	Weighted Entropy
Weak (W)	9	6	3	0.9182	$\frac{9}{15} \times 0.9182 = 0.5509$
Strong (S)	6	3	3	1.0000	$\frac{6}{15} \times 1.0000 = 0.4000$
Total	15	9	6		0.7672

Root Information Gain Comparison

Attribute	Information Gain Computation	Gain Value
Outlook	$0.9709 - 0.69088$	0.28002
Wind	$1.0000 - 0.76720$ (or $0.9709 - 0.7672 = 0.2037$)	0.23280
Humidity	$0.9544 - 0.78500$ (or $0.9709 - 0.7850 = 0.1859$)	0.16940
Temperature	$0.9709 - 0.90720$	0.06370

Ranking: Outlook > Wind > Humidity > Temp. Therefore, **Outlook** is selected as the **Root Node**.

- When Outlook = Overcast: Count = 4, all 4 Yes $\implies$ **Pure Leaf: Predict Yes.**

Partitioning the Rain Branch ($n = 5$)

The Rain subset contains 5 observations (Days 4, 5, 6, 10, 14) with 3 Yes and 2 No ($H(S_{\text{Rain}}) = 0.9709$):

Attribute & Value	Count	Yes	No	Entropy	Weighted Entropy
Temperature					
Mild	3	2	1	0.9182	$\frac{3}{5} \times 0.9182 = 0.5509$
Cool	2	1	1	1.0000	$\frac{2}{5} \times 1.0000 = 0.4000$
Weighted Sum					0.9509
Humidity					
High	2	1	1	1.0000	$\frac{2}{5} \times 1.0000 = 0.4000$
Normal	3	2	1	0.9182	$\frac{3}{5} \times 0.9182 = 0.55092$
Weighted Sum					0.95092
Wind					
Weak	3	3	0	0	$\frac{3}{5} \times 0 = 0$
Strong	2	0	2	0	$\frac{2}{5} \times 0 = 0$
Weighted Sum					0

Information Gain under Rain:

$$IG(\text{Temperature}) = 0.9709 - 0.9509 = 0.0200$$

$$IG(\text{Humidity}) = 0.9709 - 0.95092 = 0.01998$$

$$IG(\text{Wind}) = 0.9709 - 0.0000 = \mathbf{0.9709}$$

Ranking: Wind > Temp > Humidity. Under Rain, we split on **Wind**:

- Wind = Weak $\implies$ **Yes**
- Wind = Strong $\implies$ **No**

Partitioning the Sunny Branch ($n = 6$)

The Sunny subset contains 6 observations (Days 1, 2, 8, 9, 11, 15) with 2 Yes and 4 No ($H(S_{\text{Sunny}}) = 0.9182$):

Attribute & Value	Count	Yes	No	Entropy	Weighted Entropy
Temperature					
Hot	3	0	3	0	$\frac{3}{6} \times 0 = 0$
Mild	2	1	1	1.0000	$\frac{2}{6} \times 1.0000 = 0.3333$
Cool	1	1	0	0	$\frac{1}{6} \times 0 = 0$
Weighted Sum					0.3333
Humidity					
High	4	0	4	0	$\frac{4}{6} \times 0 = 0$
Normal	2	2	0	0	$\frac{2}{6} \times 0 = 0$
Weighted Sum					0
Wind					
Weak	4	1	3	0 (or 0.811)	$\frac{4}{6} \times 0 = 0$
Strong	2	1	1	1.0000	$\frac{2}{6} \times 1.0000 = 0.3333$
Weighted Sum					0.3333

Information Gain under Sunny:

$$IG(\text{Temperature}) = 0.9182 - 0.3333 = 0.5849$$

$$IG(\text{Humidity}) = 0.9182 - 0.0000 = \mathbf{0.9182}$$

$$IG(\text{Wind}) = 0.9182 - 0.3333 = 0.5849$$

Ranking: Humidity > Temp = Wind. Under Sunny, we split on **Humidity**:

- Humidity = High $\Rightarrow$ **No**
- Humidity = Normal $\Rightarrow$ **Yes**

Final Play Tennis Decision Tree




Figure 5.2: Final Decision Tree Structure for Play Tennis Classification.

5.4 Solved Problem 5.2: Credit Loan Approval Decision Tree

Example 5.4.1: Decision Tree Induction for Loan Approval

Given a dataset of 15 loan applicants across features Location, Income, Credit Score, and Employment, build the optimal decision tree using Entropy and Information Gain.

App	Location	Income	Credit Score	Employment	Approved?
1	Urban (U)	High (H)	Poor (P)	Part-time (PT)	No (N)
2	Urban (U)	High (H)	Poor (P)	Full-time (FT)	No (N)
3	Suburban (S)	High (H)	Poor (P)	Part-time (PT)	Yes (Y)
4	Rural (R)	Medium (M)	Poor (P)	Part-time (PT)	Yes (Y)
5	Rural (R)	Low (L)	Good (G)	Part-time (PT)	Yes (Y)
6	Rural (R)	Low (L)	Good (G)	Full-time (FT)	No (N)
7	Suburban (S)	Low (L)	Good (G)	Full-time (FT)	Yes (Y)
8	Urban (U)	Medium (M)	Poor (P)	Part-time (PT)	No (N)
9	Urban (U)	Low (L)	Good (G)	Part-time (PT)	Yes (Y)
10	Rural (R)	Medium (M)	Good (G)	Part-time (PT)	Yes (Y)
11	Urban (U)	Medium (M)	Good (G)	Full-time (FT)	Yes (Y)
12	Suburban (S)	Medium (M)	Poor (P)	Full-time (FT)	Yes (Y)
13	Suburban (S)	High (H)	Good (G)	Part-time (PT)	Yes (Y)
14	Rural (R)	Medium (M)	Poor (P)	Full-time (FT)	No (N)
15	Urban (U)	High (H)	Poor (P)	Part-time (PT)	No (N)

Overall Dataset Statistics and Entropy

The target variable Approved? contains 9 Yes and 6 No out of 15 applications:

$$P_Y = \frac{9}{15}, \quad P_N = \frac{6}{15} \quad (5.9)$$

$$H_{\text{Approved}} = -\left(\frac{9}{15}\right) \log_2 \left(\frac{9}{15}\right) - \left(\frac{6}{15}\right) \log_2 \left(\frac{6}{15}\right) = 0.9709 \quad (5.10)$$

Root Attribute Evaluation

1. Location ($n = 15$):

Location	Count	Yes	No	Entropy (H_i)	Weighted Entropy
Urban	6	2	4	0.9182	$\frac{6}{15} \times 0.9182 = 0.36728$
Suburban	4	4	0	0	$\frac{4}{15} \times 0 = 0$
Rural	5	3	2	0.9709	$\frac{5}{15} \times 0.9709 = 0.32363$
Total	15	9	6		0.6909

2. Income ($n = 15$):

Income	Count	Yes	No	Entropy (H_i)	Weighted Entropy
High	5	2	3	0.9709	$\frac{5}{15} \times 0.9709 = 0.32363$
Medium	6	4	2	0.9182	$\frac{6}{15} \times 0.9182 = 0.36728$
Low	4	3	1	0.8112	$\frac{4}{15} \times 0.8112 = 0.21632$
Total	15	9	6		0.9072

3. Credit Score ($n = 15$):

Credit Score	Count	Yes	No	Entropy (H_i)	Weighted Entropy
Good	7	6	1	0.5916	$\frac{7}{15} \times 0.5916 = 0.27608$
Poor	8	3	5	0.9544	$\frac{8}{15} \times 0.9544 = 0.50901$
Total	15	9	6		0.7850

4. Employment ($n = 15$):

Employment	Count	Yes	No	Entropy (H_i)	Weighted Entropy
Full-time	6	3	3	1.0000	$\frac{6}{15} \times 1.0000 = 0.4000$
Part-time	9	6	3	0.9182	$\frac{9}{15} \times 0.9182 = 0.5509$
Total	15	9	6		0.9509

Root Information Gain Comparison

Root Attribute	Information Gain	Gain Value
Location	$0.9709 - 0.6909$	0.2800
Credit Score	$0.9709 - 0.7850$	0.1859
Income	$0.9709 - 0.9072$	0.0637
Employment	$0.9709 - 0.9509$	0.0200

Ranking: Location > Credit Score > Income > Employment. Thus, **Location** is selected as the **Root Node**.

- When Location = Suburban: Count = 4, all 4 Yes $\implies$ **Pure Leaf: Predict Yes.**

Partitioning the Urban Branch ($n = 6$)

The Urban subset contains 6 observations (Apps 1, 2, 8, 9, 11, 15) with 2 Yes and 4 No ($H(S_{\text{Urban}}) = 0.9182$):

Attribute & Value	Count	Yes	No	Entropy	Weighted Entropy
Income					
High	3	0	3	0	$\frac{3}{6} \times 0 = 0$
Medium	2	1	1	1.0000	$\frac{2}{6} \times 1 = 0.3333$
Low	1	1	0	0	$\frac{1}{6} \times 0 = 0$
Weighted Sum					0.3333
Credit Score					
Good	2	2	0	0	$\frac{2}{6} \times 0 = 0$
Poor	4	0	4	0	$\frac{4}{6} \times 0 = 0$
Weighted Sum					0
Employment					
Full-time	2	1	1	1.0000	$\frac{2}{6} \times 1 = 0.3333$
Part-time	4	1	3	0.8112	$\frac{4}{6} \times 0.8112 = 0.5408$
Weighted Sum					0.8741

Information Gain under Urban:

$$IG(\text{Income}) = 0.9182 - 0.3333 = 0.5849$$

$$IG(\text{Credit Score}) = 0.9182 - 0.0000 = \mathbf{0.9182}$$

$$IG(\text{Employment}) = 0.9182 - 0.8741 = 0.0441$$

Ranking: Credit Score > Income > Employment. Under Urban, we split on **Credit Score**:

- Credit Score = Good $\Rightarrow$ **Yes**
- Credit Score = Poor $\Rightarrow$ **No**

Partitioning the Rural Branch ($n = 5$)

The Rural subset contains 5 observations (Apps 4, 5, 6, 10, 14) with 3 Yes and 2 No ($H(S_{\text{Rural}}) = 0.9709$):

Attribute & Value	Count	Yes	No	Entropy	Weighted Entropy
Income					
Medium	3	2	1	0.9182	$\frac{3}{5} \times 0.9182 = 0.55092$
Low	2	1	1	1.0000	$\frac{2}{5} \times 1 = 0.4000$
Weighted Sum					0.95092
Credit Score					
Good	3	2	1	0.8899	$\frac{3}{5} \times 0.8899 = 0.53394$
Poor	2	1	1	1.0000	$\frac{2}{5} \times 1 = 0.4000$
Weighted Sum					0.93394
Employment					
Full-time	2	0	2	0	$\frac{2}{5} \times 0 = 0$
Part-time	3	3	0	0	$\frac{3}{5} \times 0 = 0$
Weighted Sum					0

Information Gain under Rural:

$$IG(\text{Income}) = 0.9709 - 0.95092 = 0.01998$$

$$IG(\text{Credit Score}) = 0.9709 - 0.93394 = 0.03696$$

$$IG(\text{Employment}) = 0.9709 - 0.0000 = \mathbf{0.9709}$$

Ranking: Employment > Credit Score > Income. Under Rural, we split on **Employment**:

- Employment = Full-time $\Rightarrow$ **No**
- Employment = Part-time $\Rightarrow$ **Yes**

Final Loan Approval Decision Tree




Figure 5.3: Final Decision Tree Structure for Credit Loan Approval.

5.5 Solved Problem 5.3: Candidate Job Offer Decision Tree

Example 5.5.1: Decision Tree Induction for Candidate Job Offers

A human resources department assesses 14 candidates across Experience, Salary Requirement, Certification Status, and Interview Rating to predict whether a Job Offer is extended.

Candidate	Experience	Salary	Certified	Rating	Job Offer
1	Junior (J)	High (H)	No (N)	Fair (F)	No (N)
2	Junior (J)	High (H)	No (N)	Excellent (E)	No (N)
3	Mid (M)	High (H)	No (N)	Fair (F)	Yes (Y)
4	Senior (S)	Medium (M)	No (N)	Fair (F)	Yes (Y)
5	Senior (S)	Low (L)	Yes (Y)	Fair (F)	Yes (Y)
6	Senior (S)	Low (L)	Yes (Y)	Excellent (E)	No (N)
7	Mid (M)	Low (L)	Yes (Y)	Excellent (E)	Yes (Y)
8	Junior (J)	Medium (M)	No (N)	Fair (F)	No (N)
9	Junior (J)	Low (L)	Yes (Y)	Fair (F)	Yes (Y)
10	Senior (S)	Medium (M)	Yes (Y)	Fair (F)	Yes (Y)
11	Junior (J)	Medium (M)	Yes (Y)	Excellent (E)	Yes (Y)
12	Mid (M)	Medium (M)	No (N)	Excellent (E)	Yes (Y)
13	Mid (M)	High (H)	Yes (Y)	Fair (F)	Yes (Y)
14	Senior (S)	Medium (M)	No (N)	Excellent (E)	No (N)

Overall Dataset Statistics and Entropy

The target variable Job Offer contains 9 Yes and 5 No out of 14 candidates:

$$P_{\text{Yes}} = \frac{9}{14}, \quad P_{\text{No}} = \frac{5}{14} \quad (5.11)$$

$$H_S = - \left(\frac{9}{14} \right) \log_2 \left(\frac{9}{14} \right) - \left(\frac{5}{14} \right) \log_2 \left(\frac{5}{14} \right) = 0.9402 \quad (5.12)$$

Root Attribute Evaluation

1. Experience ($n = 14$):

Experience	Count	Yes	No	Entropy (H_i)	Weighted Entropy
Junior (J)	5	2	3	0.9709	$\frac{5}{14} \times 0.9709 = 0.34675$
Mid (M)	4	4	0	0	$\frac{4}{14} \times 0 = 0$
Senior (S)	5	3	2	0.9709	$\frac{5}{14} \times 0.9709 = 0.34675$
Total	14	9	5		0.6935

2. Salary ($n = 14$):

Salary	Count	Yes	No	Entropy (H_i)	Weighted Entropy
High (H)	4	2	2	1.0000	$\frac{4}{14} \times 1 = 0.2857$
Medium (M)	6	4	2	0.9182	$\frac{6}{14} \times 0.9182 = 0.3935$
Low (L)	4	3	1	0.8112	$\frac{4}{14} \times 0.8112 = 0.2318$
Total	14	9	5		0.9110

3. Certified ($n = 14$):

Certified	Count	Yes	No	Entropy (H_i)	Weighted Entropy
Yes (Y)	7	6	1	0.5916	$\frac{7}{14} \times 0.5916 = 0.2958$
No (N)	7	3	4	0.9852	$\frac{7}{14} \times 0.9852 = 0.4926$
Total	14	9	5		0.7884

4. Rating ($n = 14$):

Rating	Count	Yes	No	Entropy (H_i)	Weighted Entropy
Fair (F)	8	6	2	0.8112	$\frac{8}{14} \times 0.8112 = 0.4635$
Excellent (E)	6	3	3	1.0000	$\frac{6}{14} \times 1 = 0.4286$
Total	14	9	5		0.8921

Root Information Gain Comparison

Root Attribute	Information Gain	Gain Value
Experience	$0.9402 - 0.6935$	0.2467
Certified	$0.9402 - 0.7884$	0.1518
Rating	$0.9402 - 0.8921$	0.0481
Salary	$0.9402 - 0.9110$	0.0292

Ranking: Experience > Certified > Rating > Salary. Thus, **Experience** is chosen as the **Root Node**.

- When Experience = Mid: Count = 4, all 4 Yes $\implies$ **Pure Leaf: Predict Yes.**

Partitioning the Junior Branch ($n = 5$)

The Junior subset contains 5 observations (Candidates 1, 2, 8, 9, 11) with 2 Yes and 3 No ($H(S_{\text{Junior}}) = 0.9709$):

Attribute & Value	Count	Yes	No	Entropy	Weighted Entropy
Certified					
Yes	2	2	0	0	$\frac{2}{5} \times 0 = 0$
No	3	0	3	0	$\frac{3}{5} \times 0 = 0$
Weighted Sum					0
Salary					
High	2	0	2	0	$\frac{2}{5} \times 0 = 0$
Medium	2	1	1	1.0000	$\frac{2}{5} \times 1 = 0.4000$
Low	1	1	0	0	$\frac{1}{5} \times 0 = 0$
Weighted Sum					0.4000
Rating					
Fair	3	1	2	0.9182	$\frac{3}{5} \times 0.9182 = 0.55092$
Excellent	2	1	1	1.0000	$\frac{2}{5} \times 1 = 0.4000$
Weighted Sum					0.95092

Information Gain under Junior:

$$IG(\text{Certified}) = 0.9709 - 0.0000 = \mathbf{0.9709}$$

$$IG(\text{Salary}) = 0.9709 - 0.4000 = 0.5709$$

$$IG(\text{Rating}) = 0.9709 - 0.95092 = 0.01998$$

Ranking: Certified > Salary > Rating. Under Junior, we split on **Certified**:

- Certified = Yes $\implies$ **Yes**
- Certified = No $\implies$ **No**

Partitioning the Senior Branch ($n = 5$)

The Senior subset contains 5 observations (Candidates 4, 5, 6, 10, 14) with 3 Yes and 2 No ($H(S_{\text{Senior}}) = 0.9709$):

Attribute & Value	Count	Yes	No	Entropy	Weighted Entropy
Certified					
Yes	3	2	1	0.9182	$\frac{3}{5} \times 0.9182 = 0.55092$
No	2	1	1	1.0000	$\frac{2}{5} \times 1 = 0.4000$
Weighted Sum				0.95092	
Salary					
Medium	3	2	1	0.9182	$\frac{3}{5} \times 0.9182 = 0.55092$
Low	2	1	1	1.0000	$\frac{2}{5} \times 1 = 0.4000$
Weighted Sum				0.95092	
Rating					
Fair	3	3	0	0	$\frac{3}{5} \times 0 = 0$
Excellent	2	0	2	0	$\frac{2}{5} \times 0 = 0$
Weighted Sum				0	

Information Gain under Senior:

$$IG(\text{Certified}) = 0.9709 - 0.95092 = 0.01998$$

$$IG(\text{Salary}) = 0.9709 - 0.95092 = 0.01998$$

$$IG(\text{Rating}) = 0.9709 - 0.0000 = \mathbf{0.9709}$$

Ranking: Rating > Certified = Salary. Under Senior, we split on **Rating**:

- Rating = Fair $\Rightarrow$ **Yes**
- Rating = Excellent $\Rightarrow$ **No**

Final Candidate Job Offer Decision Tree




Figure 5.4: Final Decision Tree Structure for Candidate Job Offer Acceptance.

Chapter 6

Support Vector Machines

6.1 Core Principles and Mathematical Foundations

Support Vector Machine (SVM) is a powerful supervised learning algorithm primarily utilized for classification tasks, though it also extends to regression problems via Support Vector Regression (SVR). The objective of an SVM is to identify the optimal separating hyperplane that maximizes the geometric margin between distinct classes.

6.1.1 Fundamental Definitions and Geometry

- **Optimal Hyperplane:** A linear decision boundary that separates classes while maximizing the perpendicular distance to the closest training observations.
- **Support Vectors:** The subset of training data points that lie directly on the margin boundaries. These critical points alone determine the exact position and orientation of the separating hyperplane; all other training points can be removed or relocated without altering the solution.
- **Decision Boundary:** The separating hyperplane where the decision function evaluates to zero:

$$w^T x + b = 0 \quad (6.1)$$

- **Margin Boundaries:** The parallel hyperplanes that delimit the empty slab separating the two classes:

$$\text{Positive Margin Boundary: } w^T x + b = +1 \quad (6.2)$$

$$\text{Negative Margin Boundary: } w^T x + b = -1 \quad (6.3)$$

- **Margin Width (M):** The total geometric distance between the positive and negative margin boundaries:

$$M = \frac{2}{\|w\|} \quad (6.4)$$

where $\|w\| = \sqrt{w_1^2 + w_2^2 + \dots + w_p^2}$ is the Euclidean norm of the weight vector.

- **Perpendicular Distance of a Point:** The geometric distance of any arbitrary observation x_i to the decision boundary:

$$d_i = \frac{|w^T x_i + b|}{\|w\|} = \frac{|f(x_i)|}{\|w\|} \quad (6.5)$$

6.1.2 Functional Margin and Hard-Margin Condition

The **functional margin** of an observation (x_i, y_i) with label $y_i \in \{-1, +1\}$ is defined as:

$$\hat{\gamma}_i = y_i(w^T x_i + b) \quad (6.6)$$

For a correctly scaled canonical hard-margin SVM, all training observations must satisfy the strict separability constraint:

$$y_i(w^T x_i + b) \geq 1, \quad \forall i = 1, 2, \dots, n \quad (6.7)$$

where equality $y_i(w^T x_i + b) = 1$ holds if and only if x_i is a support vector.

6.2 Boundary Identification and Analytical Solution Techniques

To solve linear SVM problems analytically for two-dimensional data, we identify whether the classes are separated by a horizontal, vertical, or diagonal boundary:

6.2.1 Horizontal Separation ($x_2 = c$)

A dataset exhibits horizontal separation if the second coordinate (x_2) of all points in one class is strictly greater than the second coordinate of all points in the opposite class:

$$\min\{x_2 \in \text{Class } +1\} > \max\{x_2 \in \text{Class } -1\} \quad (6.8)$$

- Let $p = \min\{x_2 \in \text{Class } +1\}$ denote the smallest x_2 value on the positive side.
- Let $n = \max\{x_2 \in \text{Class } -1\}$ denote the largest x_2 value on the negative side.
- The optimal maximum margin boundary passes through the midpoint:

$$c = \frac{p+n}{2} \implies x_2 = c \quad (6.9)$$

- Setting $w = [0, w_2]^T$, we solve the canonical equations using the closest boundary points:

$$w_2 p + b = 1 \quad (6.10)$$

$$w_2 n + b = -1 \quad (6.11)$$

Subtracting the equations yields:

$$w_2(p - n) = 2 \implies w_2 = \frac{2}{p - n} \quad (6.12)$$

Substituting w_2 into either equation determines the intercept b .

6.2.2 Vertical Separation ($x_1 = c$)

A dataset exhibits vertical separation if the first coordinate (x_1) of all points in one class is strictly greater than the first coordinate of all points in the opposite class:

$$\min\{x_1 \in \text{Class } +1\} > \max\{x_1 \in \text{Class } -1\} \quad (6.13)$$

- Let $p = \min\{x_1 \in \text{Class } +1\}$ and $n = \max\{x_1 \in \text{Class } -1\}$.
- The maximum margin boundary is given by:

$$c = \frac{p+n}{2} \implies x_1 = c \quad (6.14)$$

- Setting $w = [w_1, 0]^T$, we solve:

$$w_1 p + b = 1, \quad w_1 n + b = -1 \implies w_1 = \frac{2}{p - n} \quad (6.15)$$

6.2.3 Diagonal Separation ($x_2 - x_1 = c$ or $x_2 + x_1 = c$)

When neither coordinate alone separates the classes, evaluate linear combinations such as $x_2 - x_1$ or $x_2 + x_1$. If the resulting combined score is strictly positive for one class and strictly negative for the other, the optimal decision boundary is diagonal:

$$x_2 - x_1 = c \quad \text{or} \quad x_2 + x_1 = c \quad (6.16)$$

6.3 Solved Problem 6.1: Horizontal Decision Boundary

Example 6.3.1: SVM Induction for Horizontal Separation

Given a 12-point dataset partitioned into Class +1 and Class -1:

- **Class +1:** (1, 7), (2, 8), (3, 7), (4, 9), (5, 8), (6, 7)
- **Class -1:** (1, 3), (2, 2), (3, 3), (4, 1), (5, 2), (6, 3)

Determine the maximum margin decision boundary, compute the correctly scaled parameters w and b , identify the support vectors, verify the hard-margin condition, compute the margin width and sample distances, and classify a new query observation (3, 6).

Step 1: Identify Separating Orientation and Midpoint

Inspecting the x_2 coordinates:

$$p = \min\{x_2 \in \text{Class } +1\} = 7, \quad n = \max\{x_2 \in \text{Class } -1\} = 3 \quad (6.17)$$

Since all Class +1 points have $x_2 \geq 7$ and all Class -1 points have $x_2 \leq 3$, the classes are separated horizontally:

$$7 > c > 3 \quad (6.18)$$

The maximum margin decision boundary lies at the midpoint:

$$c = \frac{p+n}{2} = \frac{7+3}{2} = 5 \implies x_2 = 5 \quad (6.19)$$

Step 2: Solve Correctly Scaled w and b

Since the boundary is horizontal, x_1 has no effect, so $w = [0, w_2]^T$. The decision function is $f(x) = w_2 x_2 + b$. Setting up the margin boundary equations:

$$w_2(7) + b = 1 \quad (6.20)$$

$$w_2(3) + b = -1 \quad (6.21)$$

Subtracting Equation (6.21) from (6.20):

$$4w_2 = 2 \implies w_2 = 0.5 \quad (6.22)$$

Substituting $w_2 = 0.5$ into Equation (6.20):

$$0.5(7) + b = 1 \implies 3.5 + b = 1 \implies b = -2.5 \quad (6.23)$$

Thus, the canonical parameters are:

$$w = \begin{bmatrix} 0 \\ 0.5 \end{bmatrix}, \quad b = -2.5 \quad (6.24)$$

and the decision function is:

$$f(x) = 0 \cdot x_1 + 0.5x_2 - 2.5 = 0.5x_2 - 2.5 \quad (6.25)$$

Step 3: Identify Support Vectors

The support vectors are the data points satisfying $f(x) = \pm 1$:

- **Positive Support Vectors** ($x_2 = 7$): (1, 7), (3, 7), (6, 7)
- **Negative Support Vectors** ($x_2 = 3$): (1, 3), (3, 3), (6, 3)

Step 4: Verify Hard-Margin Constraint $y_i(w^T x_i + b) \geq 1$

Class +1 ($y_i = +1$)		Class -1 ($y_i = -1$)	
Point (x_1, x_2)	$y_i f(x)$	Point (x_1, x_2)	$y_i f(x)$
(1, 7)	$+1 \cdot (0.5(7) - 2.5) = 1.0$	(1, 3)	$-1 \cdot (0.5(3) - 2.5) = 1.0$
(2, 8)	$+1 \cdot (0.5(8) - 2.5) = 1.5$	(2, 2)	$-1 \cdot (0.5(2) - 2.5) = 1.5$
(3, 7)	$+1 \cdot (0.5(7) - 2.5) = 1.0$	(3, 3)	$-1 \cdot (0.5(3) - 2.5) = 1.0$
(4, 9)	$+1 \cdot (0.5(9) - 2.5) = 2.0$	(4, 1)	$-1 \cdot (0.5(1) - 2.5) = 2.0$
(5, 8)	$+1 \cdot (0.5(8) - 2.5) = 1.5$	(5, 2)	$-1 \cdot (0.5(2) - 2.5) = 1.5$
(6, 7)	$+1 \cdot (0.5(7) - 2.5) = 1.0$	(6, 3)	$-1 \cdot (0.5(3) - 2.5) = 1.0$

All points satisfy $y_i(w^T x_i + b) \geq 1$ strictly, confirming hard-margin validity.

Step 5: Margin Width and Geometric Distances

$$\|w\| = \sqrt{0^2 + 0.5^2} = 0.5 \quad (6.26)$$

$$M = \frac{2}{\|w\|} = \frac{2}{0.5} = 4 \quad (6.27)$$

The perpendicular distance of each point to the boundary $x_2 = 5$ is $d_i = \frac{|0.5x_2 - 2.5|}{0.5} = |x_2 - 5|$:

Class +1		Class -1	
Point	Distance d_i	Point	Distance d_i
(1, 7)	2	(1, 3)	2
(2, 8)	3	(2, 2)	3
(3, 7)	2	(3, 3)	2
(4, 9)	4	(4, 1)	4
(5, 8)	3	(5, 2)	3
(6, 7)	2	(6, 3)	2

Step 6: Classify Query Observation (3, 6)

Evaluating the decision function at $x = (3, 6)$:

$$f(3, 6) = 0.5(6) - 2.5 = 3.0 - 2.5 = 0.5 \quad (6.28)$$

Since $f(3, 6) = 0.5 \geq 0$, the query point is assigned to **Class +1**:

$$\hat{y} = +1 \quad (6.29)$$




Figure 6.1: Horizontal SVM Maximum Margin Boundary at $x_2 = 5$ with Support Vectors and Query Observation (3, 6).

6.4 Solved Problem 6.2: Vertical Decision Boundary

Example 6.4.1: SVM Induction for Vertical Separation

Given the 12-point dataset:

- **Class +1:** (7, 1), (8, 2), (7, 3), (9, 2), (8, 1), (7, 2)
- **Class -1:** (3, 1), (2, 2), (3, 3), (1, 2), (2, 1), (3, 2)

Determine the maximum margin decision boundary, solve for w and b , identify the support vectors, verify the hard-margin condition, compute the margin and distances, and classify the query point (5, 2).

Step 1: Identify Separating Orientation and Midpoint

Inspecting the x_1 coordinates:

$$p = \min\{x_1 \in \text{Class } +1\} = 7, \quad n = \max\{x_1 \in \text{Class } -1\} = 3 \quad (6.30)$$

Since all Class +1 points have $x_1 \geq 7$ and all Class -1 points have $x_1 \leq 3$, the classes are separated vertically:

$$7 > c > 3 \implies c = \frac{7+3}{2} = 5 \implies x_1 = 5 \quad (6.31)$$

Step 2: Solve Correctly Scaled w and b

Since the boundary is vertical, $w = [w_1, 0]^T$ and $f(x) = w_1 x_1 + b$:

$$w_1(7) + b = 1 \quad (6.32)$$

$$w_1(3) + b = -1 \quad (6.33)$$

Subtracting yields $4w_1 = 2 \implies w_1 = 0.5$, and substituting gives $0.5(7) + b = 1 \implies b = -2.5$.

$$w = \begin{bmatrix} 0.5 \\ 0 \end{bmatrix}, \quad b = -2.5 \quad (6.34)$$

The decision function is:

$$f(x) = 0.5x_1 - 2.5 \quad (6.35)$$

Step 3: Support Vectors

- **Positive Support Vectors ($x_1 = 7$):** $(7, 1)$, $(7, 3)$, $(7, 2)$
- **Negative Support Vectors ($x_1 = 3$):** $(3, 1)$, $(3, 3)$, $(3, 2)$

Step 4: Verify Hard-Margin Constraint $y_i(w^T x_i + b) \geq 1$

Class +1 ($y_i = +1$)		Class -1 ($y_i = -1$)	
Point (x_1, x_2)	$y_i f(x)$	Point (x_1, x_2)	$y_i f(x)$
$(7, 1)$	$+1 \cdot (0.5(7) - 2.5) = 1.0$	$(3, 1)$	$-1 \cdot (0.5(3) - 2.5) = 1.0$
$(8, 2)$	$+1 \cdot (0.5(8) - 2.5) = 1.5$	$(2, 2)$	$-1 \cdot (0.5(2) - 2.5) = 1.5$
$(7, 3)$	$+1 \cdot (0.5(7) - 2.5) = 1.0$	$(3, 3)$	$-1 \cdot (0.5(3) - 2.5) = 1.0$
$(9, 2)$	$+1 \cdot (0.5(9) - 2.5) = 2.0$	$(1, 2)$	$-1 \cdot (0.5(1) - 2.5) = 2.0$
$(8, 1)$	$+1 \cdot (0.5(8) - 2.5) = 1.5$	$(2, 1)$	$-1 \cdot (0.5(2) - 2.5) = 1.5$
$(7, 2)$	$+1 \cdot (0.5(7) - 2.5) = 1.0$	$(3, 2)$	$-1 \cdot (0.5(3) - 2.5) = 1.0$

Step 5: Margin Width and Geometric Distances

$$\|w\| = \sqrt{0.5^2 + 0^2} = 0.5 \implies M = \frac{2}{0.5} = 4 \quad (6.36)$$

Distances $d_i = \frac{|0.5x_1 - 2.5|}{0.5} = |x_1 - 5|$:

Class +1		Class -1	
Point	Distance d_i	Point	Distance d_i
$(7, 1)$	2	$(3, 1)$	2
$(8, 2)$	3	$(2, 2)$	3
$(7, 3)$	2	$(3, 3)$	2
$(9, 2)$	4	$(1, 2)$	4
$(8, 1)$	3	$(2, 1)$	3
$(7, 2)$	2	$(3, 2)$	2

Step 6: Classify Query Observation $(5, 2)$

Evaluating at $x = (5, 2)$:

$$f(5, 2) = 0.5(5) - 2.5 = 2.5 - 2.5 = 0 \quad (6.37)$$

Boundary Indeterminacy: Since $f(5, 2) = 0$, the query point lies exactly on the decision boundary. Consequently, it cannot be strictly classified into either Class +1 or Class -1 without a tie-breaking rule.




Figure 6.2: Vertical SVM Maximum Margin Boundary at $x_1 = 5$ with Support Vectors and Query Observation $(5, 2)$.

6.5 Solved Problem 6.3: Diagonal Decision Boundary

Example 6.5.1: SVM Induction for Diagonal Separation

Given a 12-point dataset:

- **Class +1:** $(1, 3), (1, 4), (2, 4), (2, 5), (3, 5), (3, 6)$
- **Class -1:** $(3, 1), (4, 1), (4, 2), (5, 2), (5, 3), (6, 3)$

Find the optimal separating hyperplane, determine correctly scaled parameters w and b , identify support vectors, verify hard-margin conditions, compute margin and distances, and classify query observation $(3, 4)$.

Step 1: Diagnose Separating Pattern

Evaluating single coordinate boundaries reveals that neither x_1 nor x_2 separates the data independently. Testing the difference $x_2 - x_1$:

Class +1		Class -1	
Point (x_1, x_2)	$x_2 - x_1$	Point (x_1, x_2)	$x_2 - x_1$
$(1, 3)$	$3 - 1 = +2$	$(3, 1)$	$1 - 3 = -2$
$(1, 4)$	$4 - 1 = +3$	$(4, 1)$	$1 - 4 = -3$
$(2, 4)$	$4 - 2 = +2$	$(4, 2)$	$2 - 4 = -2$
$(2, 5)$	$5 - 2 = +3$	$(5, 2)$	$2 - 5 = -3$
$(3, 5)$	$5 - 3 = +2$	$(5, 3)$	$3 - 5 = -2$
$(3, 6)$	$6 - 3 = +3$	$(6, 3)$	$3 - 6 = -3$

Every point in Class +1 yields $x_2 - x_1 \geq +2$, while every point in Class -1 yields $x_2 - x_1 \leq -2$.

$$p = \min(x_2 - x_1) = +2, \quad n = \max(x_2 - x_1) = -2 \quad (6.38)$$

The maximum margin decision boundary lies at the midpoint:

$$c = \frac{p + n}{2} = \frac{+2 + (-2)}{2} = 0 \implies x_2 - x_1 = 0 \iff x_2 = x_1 \quad (6.39)$$

Step 2: Solve Correctly Scaled Parameters w and b

Rearranging the boundary equation into standard form:

$$-x_1 + x_2 + 0 = 0 \implies w = s \begin{bmatrix} -1 \\ 1 \end{bmatrix}, \quad b = 0 \quad (6.40)$$

where $s > 0$ is a positive scaling factor. Evaluating for the closest positive support point $(1, 3)$ with label $y = +1$:

$$w^T x_i = s(-1 \cdot 1 + 1 \cdot 3) = 2s \quad (6.41)$$

Setting the hard-margin condition $y_i(w^T x_i + b) = 1$:

$$+1 \cdot (2s + 0) = 1 \implies 2s = 1 \implies s = 0.5 \quad (6.42)$$

Similarly, evaluating for the closest negative support point $(3, 1)$ with label $y = -1$:

$$w^T x_i = 0.5(-1 \cdot 3 + 1 \cdot 1) = -1 \implies -1(-1 + 0) = +1 \quad (6.43)$$

Thus, the correctly scaled parameters are:

$$w = \begin{bmatrix} -0.5 \\ 0.5 \end{bmatrix}, \quad b = 0 \quad (6.44)$$

and the decision function is:

$$f(x) = -0.5x_1 + 0.5x_2 + 0 = 0.5(x_2 - x_1) \quad (6.45)$$

Step 3: Margin Boundaries and Support Vectors

$$\text{Positive Margin Boundary: } 0.5(x_2 - x_1) = +1 \implies x_2 - x_1 = +2 \quad (6.46)$$

$$\text{Negative Margin Boundary: } 0.5(x_2 - x_1) = -1 \implies x_2 - x_1 = -2 \quad (6.47)$$

The support vectors are:

- **Positive Support Vectors** ($x_2 - x_1 = 2$): $(1, 3), (2, 4), (3, 5)$
- **Negative Support Vectors** ($x_2 - x_1 = -2$): $(3, 1), (4, 2), (5, 3)$

Step 4: Verify Hard-Margin Constraint $y_i(w^T x_i + b) \geq 1$

Class +1 ($y_i = +1$)		Class -1 ($y_i = -1$)	
Point (x_1, x_2)	$y_i f(x)$	Point (x_1, x_2)	$y_i f(x)$
$(1, 3)$	$+1 \cdot 0.5(3 - 1) = 1.0$	$(3, 1)$	$-1 \cdot 0.5(1 - 3) = 1.0$
$(1, 4)$	$+1 \cdot 0.5(4 - 1) = 1.5$	$(4, 1)$	$-1 \cdot 0.5(1 - 4) = 1.5$
$(2, 4)$	$+1 \cdot 0.5(4 - 2) = 1.0$	$(4, 2)$	$-1 \cdot 0.5(2 - 4) = 1.0$
$(2, 5)$	$+1 \cdot 0.5(5 - 2) = 1.5$	$(5, 2)$	$-1 \cdot 0.5(2 - 5) = 1.5$
$(3, 5)$	$+1 \cdot 0.5(5 - 3) = 1.0$	$(5, 3)$	$-1 \cdot 0.5(3 - 5) = 1.0$
$(3, 6)$	$+1 \cdot 0.5(6 - 3) = 1.5$	$(6, 3)$	$-1 \cdot 0.5(3 - 6) = 1.5$

Hence, $y_i(w^T x_i + b) \geq 1$ holds strictly for all observations.

Step 5: Margin Width and Geometric Distances

$$\|w\| = \sqrt{(-0.5)^2 + (0.5)^2} = \sqrt{0.25 + 0.25} = \sqrt{0.5} = \frac{\sqrt{2}}{2} = \frac{1}{\sqrt{2}} \quad (6.48)$$

The total margin width is:

$$M = \frac{2}{\|w\|} = \frac{2}{\frac{1}{\sqrt{2}}} = 2\sqrt{2} \approx 2.8284 \quad (6.49)$$

The perpendicular distance from each sample point to the boundary $x_2 = x_1$ is:

$$d_i = \frac{|f(x_i)|}{\|w\|} = \frac{|0.5(x_2 - x_1)|}{1/\sqrt{2}} = \frac{|x_2 - x_1|}{\sqrt{2}} \quad (6.50)$$

Point	Class +1 Distance d_i	Point	Class -1 Distance d_i
(1, 3)	$\frac{2}{\sqrt{2}} = \sqrt{2} \approx 1.4142$	(3, 1)	$\frac{2}{\sqrt{2}} = \sqrt{2} \approx 1.4142$
(1, 4)	$\frac{3}{\sqrt{2}} = 1.5\sqrt{2} \approx 2.1213$	(4, 1)	$\frac{3}{\sqrt{2}} = 1.5\sqrt{2} \approx 2.1213$
(2, 4)	$\frac{2}{\sqrt{2}} = \sqrt{2} \approx 1.4142$	(4, 2)	$\frac{2}{\sqrt{2}} = \sqrt{2} \approx 1.4142$
(2, 5)	$\frac{3}{\sqrt{2}} = 1.5\sqrt{2} \approx 2.1213$	(5, 2)	$\frac{3}{\sqrt{2}} = 1.5\sqrt{2} \approx 2.1213$
(3, 5)	$\frac{2}{\sqrt{2}} = \sqrt{2} \approx 1.4142$	(5, 3)	$\frac{2}{\sqrt{2}} = \sqrt{2} \approx 1.4142$
(3, 6)	$\frac{3}{\sqrt{2}} = 1.5\sqrt{2} \approx 2.1213$	(6, 3)	$\frac{3}{\sqrt{2}} = 1.5\sqrt{2} \approx 2.1213$

Step 6: Classify Query Observation (3, 4)

Evaluating the decision function at $x = (3, 4)$:

$$f(3, 4) = 0.5(4 - 3) = 0.5(1) = 0.5 \quad (6.51)$$

Since $f(3, 4) = 0.5 > 0$, the query point belongs to **Class +1**:

$$\hat{y} = +1 \quad (6.52)$$




Figure 6.3: Diagonal SVM Maximum Margin Boundary $x_2 = x_1$ with Support Vectors and Query Observation (3, 4).

Comprehensive Formula Reference Sheet

Chapter 1: Machine Learning Foundations and Data Preprocessing

1. Z-Score Standardization

$$Z = \frac{x - \mu}{\sigma}$$

Where:

- x = Raw feature observation value
- μ = Mean of the feature across the dataset ($\mu = \frac{1}{n} \sum_{i=1}^n x_i$)
- σ = Standard deviation of the feature ($\sigma = \sqrt{\frac{1}{n} \sum_{i=1}^n (x_i - \mu)^2}$)
- Z = Standardized score with zero mean ($\mu_Z = 0$) and unit variance ($\sigma_Z^2 = 1$)

2. Min-Max Normalization

$$X' = \frac{X - X_{\min}}{X_{\max} - X_{\min}} \cdot (\text{new_max} - \text{new_min}) + \text{new_min}$$

Where:

- X = Original unscaled feature value
- $X_{\min}$ = Minimum observed value of the feature in the dataset
- $X_{\max}$ = Maximum observed value of the feature in the dataset
- $\text{new_min}, \text{new_max}$ = Target boundary limits (typically 0 and 1)
- X' = Rescaled feature value constrained within $[\text{new_min}, \text{new_max}]$

3. Robust Feature Scaling

$$X' = \frac{X - Q_2}{\text{IQR}} = \frac{X - \text{Median}}{Q_3 - Q_1}$$

Where:

- X = Raw feature observation value
- Q_2 = Second quartile (median / 50th percentile)
- Q_1 = First quartile (25th percentile)
- Q_3 = Third quartile (75th percentile)

- IQR = Interquartile range ($\text{IQR} = Q_3 - Q_1$)
- X' = Outlier-robust scaled feature value

4. Mean Imputation for Missing Values

$$\bar{x} = \frac{1}{n_{\text{obs}}} \sum_{i \in \mathcal{I}_{\text{obs}}} x_i, \quad x_{\text{missing}} \leftarrow \bar{x}$$

Where:

- $\mathcal{I}_{\text{obs}}$ = Index set of observed (non-missing) entries for feature x
- n_{obs} = Total number of observed values ($n_{\text{obs}} = |\mathcal{I}_{\text{obs}}|$)
- x_i = Observed feature value for sample i
- $\bar{x}$ = Arithmetic mean of observed entries imputed into missing cells

5. One-Hot Indicator Encoding

$$\mathbf{x} = \left[I(c = c_1) \quad I(c = c_2) \quad \cdots \quad I(c = c_k) \right]^T, \quad I(c = c_j) = \begin{cases} 1 & \text{if } c = c_j \\ 0 & \text{if } c \neq c_j \end{cases}$$

Where:

- c = Categorical feature taking values in unordered discrete set $\{c_1, c_2, \dots, c_k\}$
- k = Total cardinality (number of distinct categories)
- $I(\cdot)$ = Binary indicator function
- $\mathbf{x}$ = k -dimensional orthogonal binary vector representation

Chapter 2: K-Nearest Neighbors (KNN)

1. Distance Metrics in Feature Space

$$\text{Euclidean } (L_2): \quad D_E(\mathbf{x}, \mathbf{y}) = \sqrt{\sum_{i=1}^d (x_i - y_i)^2} = \|\mathbf{x} - \mathbf{y}\|_2$$

$$\text{Manhattan } (L_1): \quad D_M(\mathbf{x}, \mathbf{y}) = \sum_{i=1}^d |x_i - y_i| = \|\mathbf{x} - \mathbf{y}\|_1$$

$$\text{Minkowski } (L_p): \quad D_p(\mathbf{x}, \mathbf{y}) = \left(\sum_{i=1}^d |x_i - y_i|^p \right)^{1/p}$$

Where:

- $\mathbf{x}, \mathbf{y}$ = Coordinate vectors in d -dimensional continuous space $\mathbb{R}^d$
- x_i, y_i = Values of feature dimension i for instances $\mathbf{x}$ and $\mathbf{y}$
- d = Total number of input features
- p = Minkowski norm order parameter ($p = 1$ yields Manhattan, $p = 2$ yields Euclidean)

2. Standard Majority Voting Classification

$$\hat{Y} = \arg \max_{c \in \mathcal{C}} \sum_{i \in \mathcal{N}_k(\mathbf{x})} I(y_i = c)$$

Where:

- $\mathbf{x}$ = Unlabeled test query instance
- $\mathcal{N}_k(\mathbf{x})$ = Index set of the k nearest training instances to $\mathbf{x}$ by distance metric D
- y_i = Ground-truth class label of training instance i
- $\mathcal{C}$ = Set of all candidate target classes
- $I(y_i = c)$ = Indicator function equaling 1 if $y_i = c$, and 0 otherwise
- $\hat{Y}$ = Predicted consensus class label

3. Distance-Weighted KNN Classification

$$w_i = \frac{1}{D(\mathbf{x}, \mathbf{x}_i)}, \quad \hat{Y} = \arg \max_{c \in \mathcal{C}} \sum_{i \in \mathcal{N}_k(\mathbf{x})} w_i \cdot I(y_i = c)$$

Where:

- $D(\mathbf{x}, \mathbf{x}_i)$ = Distance from test query $\mathbf{x}$ to i -th neighbor $\mathbf{x}_i$
- w_i = Inverse distance weight assigned to neighbor i
- $\mathcal{N}_k(\mathbf{x})$ = Set of k nearest neighbors
- $\hat{Y}$ = Weighted majority predicted class label

4. KNN Continuous Regression

$$\hat{y} = \frac{1}{k} \sum_{i \in \mathcal{N}_k(\mathbf{x})} y_i$$

Where:

- $\mathbf{x}$ = Query feature vector
- y_i = Continuous target value of neighbor i
- k = Number of nearest neighbors considered
- $\hat{y}$ = Uniformly averaged predicted response value

Chapter 3: Naive Bayes Classification

1. Bayes' Theorem for Multiclass Classification

$$P(C_k | \mathbf{X}) = \frac{P(C_k) P(\mathbf{X} | C_k)}{P(\mathbf{X})} = \frac{P(C_k) P(\mathbf{X} | C_k)}{\sum_{j=1}^K P(C_j) P(\mathbf{X} | C_j)}$$

Where:

- $P(C_k | \mathbf{X})$ = Posterior probability of class C_k given feature vector $\mathbf{X}$
- $P(C_k)$ = Prior probability of class C_k
- $P(\mathbf{X} | C_k)$ = Class-conditional likelihood of feature vector $\mathbf{X}$ given class C_k
- $P(\mathbf{X})$ = Total evidence / marginal probability normalizing across all K classes

- K = Total number of distinct candidate classes

2. Class-Conditional Independence Assumption and MAP Decision Rule

$$P(\mathbf{X} | C_k) = \prod_{i=1}^d P(x_i | C_k)$$

$$\hat{C} = \arg \max_{C_k} \left[P(C_k) \prod_{i=1}^d P(x_i | C_k) \right]$$

Where:

- $\mathbf{X} = (x_1, x_2, \dots, x_d)$ = Input feature vector of dimension d
- $P(x_i | C_k)$ = Conditional likelihood of individual feature attribute x_i given class C_k
- $\hat{C}$ = Maximum A Posteriori (MAP) predicted class label

3. Maximum Likelihood Categorical Estimator and Laplace Smoothing

$$\text{MLE: } P(x_i = v | C_k) = \frac{N_{i,v,k}}{N_k}$$

$$\text{Laplace Smoothing: } P(x_i = v | C_k) = \frac{N_{i,v,k} + \alpha}{N_k + \alpha \cdot |V_i|}$$

Where:

- $N_{i,v,k}$ = Number of training instances in class C_k where feature i takes value v
- N_k = Total count of training instances in class C_k
- $|V_i|$ = Total number of possible discrete categories for feature i
- α = Additive smoothing hyperparameter ($\alpha = 1$ for standard Laplace smoothing)

Chapter 4: Logistic Regression

1. Standard Sigmoid (Logistic) Function

$$P(Y = 1 | \mathbf{X}) = \sigma(Z) = \frac{1}{1 + e^{-Z}} = \frac{e^Z}{1 + e^Z}$$

Where:

- $P(Y = 1 | \mathbf{X})$ = Predicted posterior probability of the positive class ($Y = 1$)
- Z = Linear predictor / logit argument ($Z \in (-\infty, +\infty)$)
- $\sigma(Z)$ = Sigmoid activation function mapping $\mathbb{R} \rightarrow (0, 1)$
- e = Base of the natural logarithm (≈ 2.71828)

2. Linear Predictor and Logit Transformation

$$Z = \beta_0 + \sum_{j=1}^p \beta_j X_j = \boldsymbol{\beta}^T \mathbf{X}$$

$$\text{Odds} = \frac{P}{1-P} = e^Z$$

$$\text{Logit}(P) = \ln(\text{Odds}) = \ln\left(\frac{P}{1-P}\right) = Z = \beta_0 + \sum_{j=1}^p \beta_j X_j$$

Where:

- P = Probability of the event occurring ($P = P(Y = 1 | \mathbf{X})$)
- $1 - P$ = Probability of the non-event ($P(Y = 0 | \mathbf{X})$)
- Odds = Ratio of probability of success to probability of failure
- β_0 = Intercept parameter (baseline log-odds when all $X_j = 0$)
- β_j = Regression coefficient representing change in log-odds per unit increase in X_j
- X_j = j -th explanatory feature variable

3. Odds Ratio (OR) and Percentage Change in Odds

$$\text{OR}_j = \frac{\text{Odds}(X_j + 1)}{\text{Odds}(X_j)} = e^{\beta_j}$$

$$\% \Delta \text{Odds} = (e^{\beta_j} - 1) \times 100\%$$

Where:

- OR_j = Multiplicative factor by which the odds change for a 1-unit increase in X_j
- β_j = Log-odds coefficient of feature X_j
- $\% \Delta \text{Odds}$ = Percentage increase (if positive) or decrease (if negative) in odds

4. Binary Classification Decision Rule and Cross-Entropy Cost

$$\hat{Y} = \begin{cases} 1 & \text{if } P(Y = 1 | \mathbf{X}) \geq \tau \iff Z \geq \ln\left(\frac{\tau}{1-\tau}\right) \\ 0 & \text{if } P(Y = 1 | \mathbf{X}) < \tau \end{cases} \quad (\tau = 0.5 \iff Z \geq 0)$$

$$J(\boldsymbol{\beta}) = -\frac{1}{m} \sum_{i=1}^m [y^{(i)} \ln(p^{(i)}) + (1 - y^{(i)}) \ln(1 - p^{(i)})]$$

Where:

- τ = Probability classification decision threshold (standard default $\tau = 0.5$)
- $\hat{Y}$ = Predicted binary class label ($\{0, 1\}$)
- m = Number of training examples in the dataset
- $y^{(i)}$ = Actual binary ground-truth label for instance i
- $p^{(i)}$ = Predicted model probability for instance i ($p^{(i)} = \sigma(\boldsymbol{\beta}^T \mathbf{x}^{(i)})$)
- $J(\boldsymbol{\beta})$ = Average cross-entropy (log-loss) objective function

Chapter 5: Decision Trees

1. Shannon Entropy (ID3 / C4.5)

$$H(S) = - \sum_{i=1}^c p_i \log_2(p_i)$$

$$\text{Binary Case } (c = 2): \quad H(S) = -p \log_2(p) - (1 - p) \log_2(1 - p)$$

Where:

- S = Collection of training examples at a node
- c = Number of target class categories
- p_i = Proportion of examples in S belonging to class i ($p_i = |S_i|/|S|$)
- $p, 1 - p$ = Proportions of positive and negative classes in binary classification
- $H(S)$ = Expected information entropy in bits ($0 \leq H(S) \leq \log_2 c$)

2. Gini Impurity (CART)

$$\text{Gini}(S) = 1 - \sum_{i=1}^c p_i^2$$

$$\text{Binary Case } (c = 2): \quad \text{Gini}(S) = 1 - (p^2 + (1 - p)^2) = 2p(1 - p)$$

Where:

- p_i = Probability of an item being chosen from class i
- p = Proportion of items belonging to the positive class in binary setting
- $\text{Gini}(S)$ = Measure of statistical variance of class distributions ($0 \leq \text{Gini}(S) \leq 1 - 1/c$)

3. Information Gain and Gain Ratio

$$\text{IG}(S, A) = H(S) - \sum_{v \in \text{Values}(A)} \frac{|S_v|}{|S|} H(S_v)$$

$$\text{SplitInfo}(S, A) = - \sum_{v \in \text{Values}(A)} \frac{|S_v|}{|S|} \log_2 \left(\frac{|S_v|}{|S|} \right)$$

$$\text{GainRatio}(S, A) = \frac{\text{IG}(S, A)}{\text{SplitInfo}(S, A)}$$

Where:

- $\text{IG}(S, A)$ = Expected reduction in entropy achieved by partitioning S on attribute A
- $\text{Values}(A)$ = Set of distinct branch outcomes of feature attribute A
- S_v = Subset of examples in S for which attribute A has value v
- $|S_v|/|S|$ = Weighting fraction of examples routed to child branch v
- $\text{SplitInfo}(S, A)$ = Potential information generated by splitting S into $|\text{Values}(A)|$ partitions
- $\text{GainRatio}(S, A)$ = Normalized information gain penalizing high-cardinality attributes

4. Reduction in Impurity / Gini Gain (CART)

$$\Delta \text{Gini}(S, A) = \text{Gini}(S) - \sum_{v \in \text{Values}(A)} \frac{|S_v|}{|S|} \text{Gini}(S_v)$$

Where:

- $\text{Gini}(S)$ = Impurity of the parent node before splitting
- $\text{Gini}(S_v)$ = Impurity of the child node corresponding to branch value v
- $|S_v|/|S|$ = Fraction of parent records directed to branch v
- $\Delta \text{Gini}(S, A)$ = Total impurity reduction yielded by splitting on attribute A

Chapter 6: Support Vector Machines (SVM)

1. Linear Decision Boundary Hyperplane

$$f(\mathbf{x}) = \mathbf{w}^T \mathbf{x} + b = w_1 x_1 + w_2 x_2 + \dots + w_d x_d + b = 0$$

Where:

- $\mathbf{x} = [x_1, x_2, \dots, x_d]^T$ = Input coordinate vector in $\mathbb{R}^d$
- $\mathbf{w} = [w_1, w_2, \dots, w_d]^T$ = Weight vector perpendicular (normal) to the separating hyperplane
- b = Scalar bias term governing the orthogonal displacement of the hyperplane from origin
- $f(\mathbf{x}) = 0$ = Equation of the separating decision boundary

2. Canonical Margin Hyperplanes and Total Margin Width

$$\text{Positive Margin Boundary: } \mathbf{w}^T \mathbf{x} + b = +1$$

$$\text{Negative Margin Boundary: } \mathbf{w}^T \mathbf{x} + b = -1$$

$$\text{Total Margin Width: } M = \frac{2}{\|\mathbf{w}\|} = \frac{2}{\sqrt{\sum_{j=1}^d w_j^2}}$$

Where:

- $\mathbf{w}^T \mathbf{x} + b = +1$ = Hyperplane passing through the closest positive training instances
- $\mathbf{w}^T \mathbf{x} + b = -1$ = Hyperplane passing through the closest negative training instances
- $\|\mathbf{w}\|$ = Euclidean norm (magnitude) of the weight vector ($\sqrt{\mathbf{w}^T \mathbf{w}}$)
- M = Perpendicular distance separating the positive and negative margin hyperplanes

3. Point-to-Hyperplane Perpendicular Distance

$$d_i = \frac{|\mathbf{w}^T \mathbf{x}_i + b|}{\|\mathbf{w}\|} = \frac{|f(\mathbf{x}_i)|}{\|\mathbf{w}\|}$$

Where:

- $\mathbf{x}_i$ = Coordinate vector of observation i

- $f(\mathbf{x}_i)$ = Evaluated functional value $\mathbf{w}^T \mathbf{x}_i + b$
- $\|\mathbf{w}\|$ = Euclidean norm of the normal vector
- d_i = Exact geometric perpendicular distance from sample point $\mathbf{x}_i$ to decision boundary $f(\mathbf{x}) = 0$

4. Hard-Margin Primal Feasibility Constraint

$$y_i(\mathbf{w}^T \mathbf{x}_i + b) \geq 1, \quad \forall i \in \{1, 2, \dots, n\}$$

Where:

- $y_i \in \{-1, +1\}$ = Ground-truth binary class label of training observation i
- $\mathbf{x}_i$ = Feature vector of training observation i
- n = Total number of observations in the training dataset
- $y_i(\mathbf{w}^T \mathbf{x}_i + b) = 1$ = Strict equality defining the critical support vectors lying exactly on the margin
- $y_i(\mathbf{w}^T \mathbf{x}_i + b) > 1$ = Strict inequality for interior points located beyond the margin boundaries

5. Axis-Aligned and Diagonal Separating Threshold Formulas

$$c = \frac{p + n}{2}$$

$$\text{Single Coordinate Boundary } (x_k = c): \quad w_k = \frac{2}{p - n}, \quad b = -w_k c = -\frac{p + n}{p - n}$$

$$\text{Diagonal Boundary } (x_2 - x_1 = c): \quad \mathbf{w} = s \begin{bmatrix} -1 \\ 1 \end{bmatrix}, \quad b = -s \cdot c, \quad s = \frac{2}{p - n}$$

Where:

- p = Coordinate value of the closest support vector in the positive class ($\min x_k$ or $\min(x_2 - x_1)$)
- n = Coordinate value of the closest support vector in the negative class ($\max x_k$ or $\max(x_2 - x_1)$)
- c = Optimal mid-point threshold maximizing margin width
- $w_k, \mathbf{w}$ = Canonical weight parameter(s) scaled such that $y_i(\mathbf{w}^T \mathbf{x}_i + b) = 1$ at support vectors
- b = Canonical scalar bias parameter